

Who Bears the Latency Tax? Incidence of the HFT Arms Race on a Model-Checked Limit Order Book

Atharva Patil

Xavier Institute of Engineering

Abstract

Budish, Cramton, and Shim argue that the continuous limit order book’s latency arms race taxes liquidity providers, but which participants bear it is an open empirical question, not a settled implication. We map incidence over snipe-to-quote ratio, HFT count, maker count, and demand elasticity on a C++ limit order book model-checked in TLA+, with the welfare decomposition closing to machine precision. Our headline result is within the liquidity-providing class: at the snipe-to-depth ratio measured on 86 hours of BTCUSDT-perpetual data, incidence among competing makers is redistributive — across eight latency configurations at two and three makers, one maker gains while another pays in 136 of 144 cells. Who pays is set by how the makers’ latencies are spaced, not by speed rank: with widely spaced latencies the slowest maker usually pays, in the one closely spaced set we ran it wins most often, and the pattern is identical under every assignment of latencies to agents. The class-level answer is structure-dependent — a single maker inverts the prediction and profits, through reduced pickoff losses rather than higher noise revenue, a gain that survives unit-elastic demand, while competitive supply restores it. Frequent batch auctions recover maker welfare above a minimum clearing cadence.

1. Introduction

1.1 The arms-race argument

Budish, Cramton, and Shim (2015, hereafter BCS) argue that the continuous limit order book—the dominant design of modern equity and futures markets—contains a structural flaw that has nothing to do with any individual participant’s conduct. Because the book processes messages serially in continuous time, each arrival of public information creates a fleeting, mechanical arbitrage opportunity: whichever trader reacts first can pick off the stale quotes of liquidity providers before those quotes can be revised. The competition to be first is a contest over speed rather than price, and in equilibrium it dissipates real resources into a socially wasteful arms race while transferring rent from liquidity providers (and, indirectly, the investors they serve) to the fastest “snipers.” Crucially, the resulting latency tax widens effective spreads and degrades liquidity without improving price discovery; it is a prisoner’s dilemma baked into the market’s microstructure. BCS propose frequent batch auctions—discrete-time, uniform-price clearing—as a design remedy that converts competition on speed back into competition on price, neutralizing the mechanical sniping advantage.

1.2 The incidence question, and what we find

This paper asks which participants bear the latency tax. The BCS argument establishes that the tax exists; it does not settle who pays it, and that is a separate empirical question. We map incidence over snipe-to-quote size, the number of snipers, the number of competing makers and the elasticity of uninformed demand, on a limit order book whose matching design is model-checked in TLA+ (§3.1). Two results carry the paper, and they must be read together.

- **Incidence inverts under single-maker supply (§4.1).** Mapping a 63-cell surface over snipe-to-quote size ratio and HFT count, holding liquidity supply at a single maker throughout, yields a monotone regime boundary: the maker bears the tax when snipes are large relative to quoted depth, and at small snipe sizes it gains instead, with noise traders

bearing the entire transfer. Decomposing that gain shows it is loss reduction rather than revenue capture: the maker’s noise-flow revenue *falls* when it widens, since widening sheds volume faster than it recovers edge per fill, and the gain comes from the accompanying drop in adverse-fill losses, with inventory carry supplying a third to a half of the net effect (§4.1). The ratio measured on 86 complete book-and-trade hours of the perpetual, 0.0072 at the median, falls inside that maker-gains region at every HFT count we test.

- **Competitive supply restores BCS, and within-class incidence is redistributive (§4.2).** Re-running the surface at maker counts $M \in \{1, 2, 3, 5\}$ — 441 cells, the $M = 1$ column nesting §4.1 to zero absolute difference — the maker-gains region at the tape band shrinks at $M = 2$, the smallest competitive case, to three of eighteen cells at the lowest HFT counts, and is gone at $M = 3$; elsewhere the maker’s PnL delta turns negative with its interval excluding zero. The noise-trader toll shrinks in step, so competition reinstates the incidence BCS predicts rather than merely cancelling the inversion. But that aggregate is a sum over unlike makers. Decomposing it at $n = 100$ across eight latency configurations finds one maker gaining while another pays in 136 of 144 tape-band cells. Who pays is not ordered by speed: the slowest maker pays in most cells when latencies are widely spaced, but at $\{2, 3, 4\}$ ticks — the one closely spaced set we ran — it is the most frequent winner, and unless one maker has zero latency the fastest of three never wins consistently. Permuting which agent holds which latency leaves every verdict unchanged, so this is a property of the latency configuration, not of agent order. A companion robustness experiment (§4.2) relaxes the inelastic noise-trader assumption up to unit elasticity at a single maker and leaves the single-maker inversion intact — seventeen of eighteen tape-band cells still classify as maker gains, the one loss turning indeterminate rather than reversing.

The result is therefore a boundary over two structural axes rather than a verdict on a class. Whether liquidity providers bear the latency tax depends on how many of them supply the book and on how the class is weighted; at the measured ratio, the tax is redistributed within the class, and who bears it depends on how the makers’ latencies are spaced. §§4.3–4.4 are what license that reading.

What the incidence claim rests on. Four supporting experiments establish the properties an incidence claim has to assume, each run first at a pre-registered operating point chosen for mechanism clarity and then at parameters fitted to real order flow (§§4.3–4.4):

- **The transfer is real and the accounting closes (Experiment 1, §4.3).** With latency-advantaged HFTs present we measure a positive HFT rent of 34.70 (95% CI [26.25, 44.80]) that is a near-exact zero-sum transfer from the market maker and noise traders, the welfare residual closing to machine zero (mean 2.8×10^{-12}), while liquidity gaps rise from 2.84 [2.47, 3.26] in the matched no-HFT baseline to 9.70 [9.15, 10.27] in treatment; the result holds in all nine cells of a 3×3 sensitivity-by-decay grid. Absent a decomposition that closes, any statement about who bears the tax would be an accounting artifact.
- **The mechanism carries no hidden threshold (Experiment 2, §4.3).** Rent rises with market-maker staleness at approximately 10.46 per tick, and the data do not distinguish a linear specification from a segmented one ($\Delta\text{BIC} = -0.17$, inside the range of indifference), with rent first becoming statistically distinguishable from zero at roughly two ticks. A boundary mapped across a smooth mechanism is a boundary rather than an artifact of a discontinuity.
- **Fragility rises with competition; the rent pool is fixed only by construction (Experiment 3, §4.3).** When one snipe clears the quote, total rent is flat in k (mean 34.4, range 33.70–34.93 across $k = 1 \dots 21$) — a consistency check, since one HFT wins each race — while when it does not, k enlarges the pool sevenfold by $k = 21$. Liquidity gaps grow linearly in k (fitted $\approx 2.1k + 3.7$, $R^2 = 0.997$), roughly $17\times$ by 21 HFTs, and the

growth survives switching off the maker’s depth-thinning rule; its channel is unidentified. This is what makes HFT count a real second axis of the incidence surface rather than a proxy for aggregate pressure.

- **The clearing rule drives the outcome (Experiment 4, §4.3).** Switching the engine to frequent batch-auction clearing recovers market-maker welfare monotonically over intervals ≥ 5 — about 97% of the loss by an interval of 500 (MM loss $28.5 \rightarrow 0.7$) — and this MM-recovery signal is the robust one, holding under both tie-break rules we have run, though on 50,000-tick runs its size depends on the rule. HFT-rent point estimates also fall, but only the interval-5 arm attains a rent interval clearing zero. An interval of 1 backfires (MM loss worsens to 41.3, about 137 gaps), and a slow maker alone does not produce that: a seven-tick maker still gains from a five-tick auction.
- **The environment is disciplined by real order flow (§3.6, §4.4).** Fitting to 27 days of BTCUSD-perpetual order flow and re-running all four experiments at the fitted parameters leaves every qualitative finding intact, and supplies the magnitude the operating point cannot: normalised by traded notional, HFT rent is 0.125 bp at the operating point and 4.42 bp calibrated. The calibrated figure is an upper bound, because a single latency-disadvantaged maker absorbs the whole race.

1.3 Contributions

This paper makes three contributions.

First, it treats the incidence of the latency tax as an empirical question distinct from its existence, and maps it over market structure rather than answering it with a verdict. Under single-maker supply the BCS prediction inverts at the snipe-to-depth ratio measured on the tape; a second competing maker all but reverses it back, and a third completes the reversal (§§4.1–4.2). Incidence is a function of market structure — of snipe size, sniper count, and the number of makers supplying the book — and neither BCS nor the direct message-level measurement of Aquilina, Budish, and O’Neill (2022) treats it as one.

Second, it shows that incidence within the liquidity-providing class is redistributive, so the class-level question is not well posed without a weighting: one competing maker gains while another pays in 136 of 144 cells, and who pays is set by how the makers’ latencies are spaced rather than by speed rank. The slowest maker pays in most configurations (112 of 144 cells) but is the most frequent winner in the one closely spaced set we ran (§4.2).

Third, as supporting methodology rather than a headline claim, it runs the BCS mechanism on a matching engine whose design is model-checked, and uses differential testing against that engine as an internal-validity check. The check is operative, not nominal: a Python re-implementation of the batch clearing rule, checked against the C++ source line by line, cleared a volume-tied auction at a different price the first time the two ran against the same book, and the difference moved published batch results (§3.1, Appendix A).

The negative and nuanced results — linear rather than threshold scaling, fragility that survives the depth rule, MM recovery as the clean remedy signal, and the failure of incidence to reduce to aggregate snipe pressure — are integral to the contribution.

2. Related Literature

This paper sits at the intersection of five strands of work: the market-design theory of the HFT arms race and its direct measurement, the theory of speed and exchange design, the game theory of predatory liquidity provision, the empirical and physics-of-markets literature on endogenous fragility, and the methodological tradition of agent-based modeling. We synthesize each in turn and then state our differentiator.

The arms race and its remedy. Our central reference is Budish, Cramton, and Shim (2015), who reframe latency competition not as a behaviour to be regulated but as a symptom of the continuous limit order book’s serial-processing design, and who propose frequent batch auctions as the corrective. Our work is a controlled experimental instantiation of their model: we build the latency-differentiated observation structure and submission-timing race they describe, reproduce the predicted treatment effects, and test their remedy by switching the clearing rule of the same verified engine, supplying a closed welfare decomposition their reduced-form empirics cannot directly observe. The race has since been measured directly: Aquilina, Budish, and O’Neill (2022), using message-level data that records failed race attempts alongside successful ones, find races roughly once per minute per FTSE 100 symbol, decided by microseconds, summing to a latency-arbitrage tax of about 0.4 basis points of volume. Their measurement and our experiment are two views of one transfer — they infer it from exchange messages, we observe every leg of it under a conserved accounting identity. Shkilko and Sokolov (2020) supply the closest natural experiment to our incidence question: when weather disrupts the microwave links that give the fastest traders their edge, adverse selection falls and trading costs decline. That is the class-level BCS incidence our competitive-supply arms recover (§4.2) and our single-maker arm inverts (§4.1), and their setting — many competing liquidity suppliers — is the one in which we find it. Menkveld (2016) surveys the ground between, cataloguing HFT’s double role as liquidity supplier and source of adverse selection, and noting how little of the literature settles the net welfare sign for any participant class. The two roles are documented separately. Menkveld (2013) follows a single high-frequency market maker that supplied most of the liquidity on a new venue, earning on the spread while losing on its positions; Brogaard, Hendershott, and Riordan (2014) find that HFTs’ liquidity-demanding trades impose adverse selection while their liquidity-supplying orders are adversely selected. Both place fast traders on the maker side of the book as well as the sniping side, which is where our within-class result sits. The incidence question of §4.1 sharpens exactly that unsettled sign.

Speed and exchange design. Menkveld and Zoican (2017) show that exchange speed is not neutral infrastructure: a faster engine accelerates snipers and makers alike, raising the frequency of speed duels, so lower latency can *worsen* liquidity when news arrival is high relative to liquidity-motivated trade. Our degenerate interval-1 cell, where the most aggressive clearing cadence amplifies the pathology it was meant to cure, is a discrete-clearing cousin of that comparative static. Baldauf and Mollner (2020) compare design responses — batch auctions and asymmetric speed bumps — and show both blunt the sniping externality with different effects on price discovery; we test the first experimentally, and the second is implementable in the latency scheduler without touching the matcher. Du and Zhu (2017) ask not whether discrete clearing helps but what frequency is optimal, batching longer letting more information accumulate into each price while delaying incorporation and thinning participation; Experiment 4 is the experimental counterpart of that trade-off, and what we add is a closed welfare decomposition and a fragility channel an optimal-frequency calculation does not price. Hoffmann (2014) supplies the equilibrium logic our maker executes mechanically: fast traders’ ability to revise ahead of news concentrates picking-off risk on slow limit-order submitters, who rationally quote less aggressively. Our maker’s adverse-selection widening (§3.2) is that response imposed as a rule rather than derived, which is why its parameterisation is a limitation we test directly (§5.2). Biais, Foucault, and Moinas (2015) go further and make speed itself a choice: institutions invest in fast access that lets them trade ahead of slow ones, the fast impose adverse selection on the slow, and equilibrium investment in speed exceeds the social optimum. Our within-class result is related to that fast–slow externality without reducing to it: the tax is redistributed among makers, but who pays is set by how their latencies are spaced rather than by rank — the slowest usually pays when latencies are far apart, and in the one closely spaced set we ran it wins most often (§4.2). Where they derive who is fast and what it costs in equilibrium, we impose both the latencies and the maker’s rule, and a configuration-dependent winner is precisely what an

equilibrium model would have to pin down.

Predatory versus cooperative liquidity. Carlin, Lobo, and Viswanathan (2007) provide the strategic micro-foundation for endogenous liquidity evaporation: when cooperation among strategic traders breaks down, agents race to trade ahead of a distressed counterparty and amplify price impact, producing crises that owe nothing to exogenous shocks. The BCS sniping mechanism is a latency-mediated instance of that predatory equilibrium; we differ in running an explicit agent-based experiment with a closed accounting identity rather than a stylized analytic game. Foucault, Kozhan, and Tham (2017) isolate the informational subset empirically, showing that *toxic* arbitrages — those triggered by news rather than transient imbalance — impose the adverse selection that widens spreads. Their toxic/non-toxic split is the empirical analogue of our maker’s adverse-fill classifier (§3.2), and the incidence result turns on how much the widening it triggers cuts the toxic side’s take relative to what it costs on the non-toxic side.

Adverse selection and the maker’s spread. The tradeoff our mechanism runs through predates the latency literature. Copeland and Galai (1983) model the quoted spread as the outcome of a maker’s choice between revenue collected from uninformed traders and losses to better-informed ones, so that widening is simultaneously a defence against the second and a tax on the first — which is precisely the two-legged structure the decomposition table of §4.1 measures, and their framing is why we report the legs separately rather than only their difference. Glosten and Milgrom (1985) make the adverse-selection component of the spread explicit as an information cost: even a competitive, risk-neutral, zero-inventory-cost maker must quote a strictly positive spread, because each trade carries the possibility that the counterparty knows more, and the spread is the price of that possibility. Our maker is a mechanical realization of that logic in the time domain rather than the information domain — its disadvantage is latency rather than private information — but the object it is defending against is the same, and the widening rule of §3.2 is a discrete-time analogue of their information component adjusting to the observed rate of informed trade.

Endogenous fragility and flash crashes. Kirilenko, Kyle, Samadi, and Tuzun (2017) dissect the May 2010 Flash Crash from audit-trail data and find that HFTs, while not the trigger, consumed liquidity and passed inventory among themselves at the critical moment, accelerating the decline. Experiment 3 speaks directly to that episode by asking whether such fragility can arise from rational HFT behaviour alone, with no exogenous shock; we find it grows with the number of competing HFTs, and where Kirilenko et al. perform forensics on one historical event we offer a reproducible counterfactual in which HFT count, latency, and market design are knobs. Filimonov and Sornette (2012) supply the diagnostic we borrow: the branching ratio of a self-exciting Hawkes process as a measure of reflexivity, which spikes toward criticality around fragile periods. We estimate it on simulated order flow via a Hawkes bridge, the advantage being that the data-generating process is known, so the ratio can be tied causally to HFT count and latency rather than inferred post hoc (Appendix C).

Agent-based modeling and its verifiability gap. Methodologically, our work extends the agent-based computational finance tradition (Farmer and Foley, 2009; LeBaron, 2006), which argues that bottom-up simulation of heterogeneous, boundedly rational agents can generate emergent, out-of-equilibrium phenomena—fat tails, volatility clustering, bubbles and crashes—that representative-agent equilibrium models cannot. That tradition’s well-known weaknesses are sensitivity to ad hoc behavioral assumptions, weak calibration discipline, and limited reproducibility or verifiability of the simulation machinery itself. The closest agent-based precursor to this paper is Wah and Wellman (2013), who build a tactical agent-based model of latency arbitrage across fragmented markets and find that a latency arbitrageur profits at the expense of other traders while degrading allocative efficiency — qualitatively the transfer we decompose, established a decade before us. Our departure from that work is therefore not the agent-based

method, which it had already applied to this mechanism, but the substrate the method runs on and the question we ask of its output. We inherit the tradition’s strength (emergent crash dynamics from rational behavior) while directly confronting its central weakness: rather than building on a bespoke, unverified simulator whose matching logic could itself generate the observed artifacts, we drive a TLA+-verified C++ matching engine, and we impose calibration and welfare-accounting discipline in the form of a closed, near-zero welfare residual.

Differentiator. Agent-based study of latency arbitrage is not itself new: Wah and Wellman (2013) modelled it and recovered the same qualitative transfer, so novelty of approach is not what we claim. What distinguishes this paper is, to our knowledge, the first experimental test of the BCS welfare decomposition on a *model-checked* matching engine, and the first treatment of the latency tax’s *incidence* as a question distinct from its existence. BCS is theory plus reduced-form empirics, since joined by the direct message-level measurement of Aquilina, Budish, and O’Neill (2022); the agent-based tradition, Wah and Wellman included, runs on bespoke simulators whose matching layer is a candidate source of the very effects under study. Here the entire BCS mechanism lives in a Python scheduler atop a deterministic matcher whose continuous-matching specification was model-checked with zero invariant violations (Appendix A), reached through a pybind11 bridge, so emergent arms-race dynamics, liquidity gaps, and crash behavior cannot plausibly be artifacts of the matching design (in the bounded sense of §3.1). On that substrate we decompose welfare into market-maker, noise-trader, and HFT components, show the arms race is a transfer that closes to a machine-precision zero-sum residual, trace the mechanism across latency and HFT-count sweeps (with the findings that scaling is linear rather than threshold and that fragility rises with HFT count), and confirm the BCS remedy experimentally via the batch-auction treatment, with a placebo-validated endogeneity diagnostic in the Filimonov–Sornette spirit that returns a null at our operating point. It is the joint presence of a verified engine, a closed welfare identity, and an experimental confirmation of both the pathology and its proposed cure that situates this paper relative to the literature it builds on.

3. Model and Experimental Design

This section describes the simulation apparatus. The design rests on a deliberate separation of concerns: a single model-checked C++ matching engine adjudicates every trade — continuous matching and batch-auction uncrossing alike — while a Python layer above it supplies the agents, the latency model, and the batch-auction cadence. The verified engine is the part we trust to be correct; the Python layer is the part we vary. The boundary between them is the pybind11 bridge, and it is drawn precisely where the engine’s formal guarantees end.

3.1 The matching engine, and what its verification covers

All trades in both arms — continuous matching and batch-auction uncrossing — are produced by one C++20 matching engine, reached from Python through a pybind11 bridge that is the agents’ only channel to it. The engine runs synchronously, so the fills an order produces are drained immediately after it is submitted, and it has no notion of agent latency: the entire BCS race lives in the scheduler above it (§3.3). Its circuit breaker is disabled, so that simulation control rather than an exchange halt governs the dynamics.

The engine’s matching design is model-checked in TLA+. At its deepest exhaustive configuration the continuous-matching specification explores 171,187,419 distinct reachable states with zero invariant violations, checking FIFO execution and conservation; the auction specification checks that every uncross clears at a single price, the path the batch arm runs through. Two boundaries on that verification deserve equal precision, because “model-checked” is easy to over-read. TLC’s guarantee is exhaustive only within small bounds — a handful of orders, two prices, short horizons — so it certifies the matching *design* at bounded scope rather than proving unbounded

correctness. And the object verified is the specification, not the C++ source: no refinement proof links the two, and the correspondence rests on construction and conformance testing. What the provenance buys is the elimination of design-level matching errors — the kind that could manufacture spurious gaps or transfers — within the checked scope, on the code path every trade traverses. Appendix A lists the specifications, invariants and configurations.

Differential testing against the verified engine is what makes that scope operative rather than nominal. An earlier version of the harness cleared the batch auction through a Python re-implementation of the engine’s clearing rule, checked against the C++ source line by line. Run against the same book, the two disagreed: where two prices tie on maximum executable volume, the Python rule cleared at 100 and the engine at 101. Re-routing Experiment 4 through the engine moved HFT rent in three of seven long-grid batch cells across zero in both significance and sign, and the maker’s PnL in five, while the continuous arm, which never calls the auction, returned byte-identical results — confining the change to the substituted code path. The episode also makes the tie-break rule an economic assumption the batch result depends on; Appendix A reports what survives both rules.

3.2 Agents

All agents are driven by the scheduler through one interface — each tick an agent emits orders, and it receives the fills they produce — and each holds its own cash and inventory for the welfare accounting of §3.4. Constants are in Appendix B; what matters here is each agent’s role and decision rule.

The fundamental is the exogenous true value $V(t)$, a discrete random walk in fixed-point price space. Its latency-delayed read returns the most recent V at or before $t - \ell$, and that delay is the BCS information gap separating a fast agent from a slow one; the undelayed value is used only for the maker’s adverse-fill test.

The market maker is an adverse-selection-aware two-sided quoter. On each requote tick it decays its half-spread toward base, observes V with its larger latency — so its quotes are stale — and posts a bid and ask around that stale view, skewed by inventory. Two coupled defences drive the mechanism. It widens *only* on adverse fills, meaning fills where it bought above or sold below the true current V ; noise fills sitting inside the quoted spread leave it untouched. And liquidity thins as it widens: resting size scales as base half-spread over current half-spread, so spread and depth are one coupled state rather than two independent knobs.

Each HFT is a latency-advantaged stale-quote sniper. It observes V with little or no delay, and fires an aggressive IOC when the gap between V and the resting quote exceeds its transaction cost. The race is a random submission jitter added to each order’s submit time: with several HFTs the lowest draw wins and the others find the quote already gone. Its counterparty is the slower maker, and that latency gap is the BCS arbitrage.

A noise trader is an uninformed liquidity taker; together the noise traders are the control population. Each tick it crosses the spread with fixed probability on a randomly chosen side, carrying no fundamental signal and never sniping, so it pays the half-spread and its expected PnL is negative by construction. Arrival is inelastic in spread width in every experiment through the competition arm; the elastic-demand arm relaxes exactly this, scaling the crossing probability by $(s_{\text{base}}/s)^\alpha$, with $\alpha = 0$ recovering the inelastic path bit-for-bit.

3.3 The race and the liquidity-gap signature

A latency scheduler supplies the race the engine deliberately lacks. Each tick it steps the fundamental, lets every agent emit orders, and queues each at its submission time — the current time plus the agent’s latency and, for an HFT, a random jitter. It then submits every order

whose time has arrived, in submission-time order, draining fills after each submission and routing them to the buyer and seller. The BCS submission race is exactly many agents stamping near-identical submission times, resolved in arrival order: there is no special-case race logic, only priority over arrival time, which is what makes the race an emergent consequence of the latency model rather than a scripted outcome. The scheduler records a per-tick snapshot of the book and the fundamental, and every fill.

The endogenous-instability signature of this mechanism is not a mid dislocation but a **liquidity gap**: a transient one-sided or empty book, where a thinned level is emptied before the maker requotes. An episode qualifies only if it is a maximal run of non-two-sided snapshots that recovers to a two-sided book within 50,000 μs (longer runs are structural outages and are discarded) *and* is endogenous, meaning the fundamental is flat across it: $|V_{\text{hi}} - V_{\text{lo}}| < k_{\text{fund}} \sigma \sqrt{g + 2}$ with $k_{\text{fund}} = 3$. Value-driven gaps are excluded by construction. We use this event definition because a variant maker that widened on every fill rather than only on adverse ones produced no HFT-specific divergence in mid prices: the instability this mechanism generates surfaces as gaps in depth, not crashes in price.

3.4 Metrics

All metrics are computed from the per-tick snapshots and fills, and apply unchanged to the continuous and batch arms.

Welfare decomposition and the zero-sum identity. Each agent’s PnL is its cash plus its inventory marked at the last two-sided mid (else the last fundamental value). We report the maker’s PnL, the noise traders’ summed PnL and the HFTs’ PnL; HFT rent is the HFTs’ PnL, the rent the arms race extracts; and the zero-sum residual is the sum of the three. In a closed market the residual is zero up to floating-point error. It is a bookkeeping check, not a correctness check: every fill is bilateral and every agent is marked at one price, so cash and inventory sum to zero under any correct accounting whether or not the matching is right. A non-trivial residual would indict the harness’s accounting; a zero one shows only that the accounting has no leak. Matching correctness rests instead on the model checking and the differential test of §3.1 and Appendix A. The treatment-minus-baseline welfare delta is the paper’s primary treatment effect, because it moves even when quoted spreads do not.

Units. PnL, rent and their deltas are in simulation currency — engine price units times quantity units — not US dollars, and their levels are not comparable across environments. The operating point quotes 50 units a side; the calibrated environment of §3.6 quotes about 2,061, one unit being the perp’s median trade size, which is why operating-point rents are tens of units and calibrated rents tens of thousands. Comparisons across the two use basis points of traded notional (§4.4).

Lagged Kyle’s lambda is the OLS slope of the mid change over the following five ticks on signed volume. The forward lag is deliberate: a latency-disadvantaged maker quotes stale, so a same-tick regression would score its delayed catch-up requote as *negative* impact; measuring over the next five ticks captures the catch-up as the *positive* permanent impact it actually is, removing the stale-quote sign artifact.

Liquidity gaps are counted as defined in §3.3: the endogenous-instability metric matched to the thinning mechanism — gaps, not mid crashes.

Confidence intervals are nonparametric percentile bootstraps of the mean — 2,000 resamples from a seeded generator — reported at 95%. Implementation details are in Appendix B.

3.5 Experimental design

The pre-registered operating point pairs a baseline and a treatment that share every market parameter and differ **only in HFT presence**: the **baseline** is the latency-disadvantaged, adverse-selection-aware maker plus noise traders; the **treatment** adds HFTs. Both run on the continuous scheduler through the verified engine.

It fixes 3,000 ticks, eight noise traders against three HFTs, a 100-tick base spread, per-step $\sigma = 20$, a three-tick maker latency against zero-latency snipers, and 20 seeds per arm with 2,000-resample bootstrap intervals. The full parameter set, including the agent-level constants of §3.2, is in Appendix B.

Around this operating point we run four one-dimensional sweeps, each isolating one axis of the mechanism:

1. **Calibration robustness (3×3)**. The maker’s adverse sensitivity $\times$ spread decay over $\{0.05, 0.10, 0.20\}$, confirming the result is not an artifact of the calibrated cell — HFT rent must be significantly positive and the zero-sum identity must hold in all nine cells.
2. **Latency sweep (Exp 2)**. The maker’s staleness gap from 0 to 20 ticks with HFTs at zero latency, on a tick-aligned grid (latency resolves only at tick granularity), testing whether rent rises linearly or exhibits a knee/phase transition.
3. **HFT count sweep (Exp 3)**. Number of HFTs over $\{0, 1, 2, 3, 5, 8, 13, 21\}$, separating per-HFT rent (competed away) from total rent (transfer) and from fragility (liquidity gaps).
4. **Auction-frequency sweep (Exp 4)**. Batch interval over $\{1, 5, 10, 25, 50, 100, 500\}$ plus continuous, run through a batch scheduler that parks the book in an auction state and clears through the engine’s own uncross (Appendix A), measuring how much maker welfare loss the frequent batch auction recovers and at what interval.

Each sweep varies one parameter and holds the rest at the operating point, so that any change in the outcome metrics is attributable to that single axis. We note that the no-HFT baseline liquidity-gap figure is reported per experiment from its own matched run. At $n = 100$ these agree: the matched baseline anchoring the Experiment 1 headline comparison is 2.84 gaps (95% CI [2.47, 3.26]), and the Experiment 3 $k = 0$ cell that anchors the fragility ratio reports the same 2.84.

3.6 Calibration to BTC order flow

To discipline the operating point against real microstructure, we calibrate the sim’s *environment* parameters to Binance BTCUSDT-perpetual order flow recorded by a companion collector (20-level book snapshots at $\sim 10/s$ and the aggregated-trade stream, both with nanosecond timestamps). The collector drops a few hours on most days (host sleep), which we handle by making each moment gap-safe rather than by discarding gapped days: arrival intensity accumulates each hourly file’s own observed span, so missing hours leave the denominator instead of inflating it, and realized volatility keeps only returns between *adjacent* one-second buckets, so no return straddles a gap. Days enter at 18 hours of coverage in both streams; files truncated by a mid-write crash are skipped and recorded. This is not housekeeping: a whole-day span over a 21/24 day biases the arrival rate down by roughly 12%, and demanding 24/24 coverage admits only two of the days the collector has ever produced. Under the gap-safe estimators the sample is 27 days (2026-06-24 to 2026-07-21, less 2026-07-19; 37.0M trades, 4.16M snapshots, 2.27M seconds of coverage), yielding: trade arrival 16.28/s; trade sizes p50/p90/p99 of 0.003/0.177/2.01 BTC; median spread of one exchange tick (≈ 0.016 bp); median top-of-book depth 6.18 BTC; and 1-second realized volatility of 0.650 bp/ $\sqrt{s}$.

The mapping between sim ticks and wall-clock is a modelling choice we fix at 1 tick = 100 ms:

the market maker’s 3-tick latency then means 300 ms of quote staleness — slow but plausible — while HFTs react within a tick, preserving the BCS fast/slow asymmetry at realistic crypto scales. Under that convention two environment parameters invert directly, and both fits verify by simulation. The volatility inversion is clean: at fitted σ the baseline’s simulated realized vol is 0.647 vs the target 0.650 bp/ $\sqrt{s}$. The arrival inversion is not: its closed form assumes every noise order trades, but only about half the IOCs find a quote at the operating point, so the fit iterates the arrival intensity against the *simulated* trade rate to a fixed point (converging in three steps from a closed-form 0.204 to 15.94 vs 16.28/s). The fitted values are a per-tick arrival probability of 0.427 and $\sigma = 20.6$ against the pre-registered operating point’s 0.15 and 20.

The volatility figure is worth dwelling on. Estimated over the full month the fitted σ lands within 3% of the value the experiments were pre-registered at, so on volatility the operating point is not data-adjacent but data-consistent. A two-day subsample passing a stricter 24/24 completeness filter gives 55.7 — nearly three times as large — because those two days were simply hot (1.76 bp/ $\sqrt{s}$ against the 27-day 0.650), which is the kind of sampling artifact a small clean-data-only window invites. Arrival intensity remains genuinely under-set: the fitted per-trader rate is $2.8\times$ the operating point’s, so the supporting experiments (§4.3) run on a market thinner in order flow than the perp, and §4.4 re-runs the full grid at the fitted parameters.

That re-run required re-calibrating one agent parameter, for a reason that generalizes to any transplant of a tuned agent into a calibrated environment. The maker’s widening is applied *per adverse fill event* (each adverse fill multiplies the half-spread by one plus the sensitivity), so it is not scale-free in order flow: tripling the arrival rate triples the multiplicative kicks per tick while the decay pulling the spread back to base is unchanged. Measured directly, the maker’s pickoff rate rises from 0.20 to 0.86 per tick between the operating point and calibrated flow, and at the pre-registered adverse sensitivity of 0.1 the *no-HFT* spread settles at $5.36\times$ base — the maker widens to its cap with no informed trader in the market at all. Any experiment run in that state would credit a control-loop artifact to the arms race. The stationary spread multiple obeys $m = (1 - P \ln(1 + s)/d)^{-1}$ for pickoff rate P , sensitivity s and decay d , which predicts $5.56\times$ at $P = 0.86$ against the $5.36\times$ observed and $1.24\times$ at $P = 0.20$ against $1.17\times$; solving it for the original calibration target locates the replacement value, and a re-run of the §3.5 sensitivity-decay grid under calibrated flow confirms it. We therefore set the adverse sensitivity to 0.015 for the calibrated runs, holding spread decay at its pre-registered 0.1 so that exactly one parameter moves, selected on the same criterion as the original (no-HFT spread $\approx 1.1\times$ base with positive with-HFT divergence: $1.12\times$ and $+0.08$, against $1.14\times$ and $+0.03$ for the published cell). The $6.7\times$ reduction in sensitivity tracks the $4.3\times$ rise in pickoff rate, as per-event widening implies.

Two further parameters are not pinned by the data. The snipe size is set to the calibrated quote size (2061 units), preserving the operating point’s ratio of 1.0 so that a single snipe still clears the top quote — the mechanism the race runs on; holding it at 50 instead makes a snipe take 2.4% of the book, which changes the experiment rather than rescaling it, and we report that arm separately as a robustness check. The minimum quote size is left at 1 and is inert at calibrated depth: the thinning floor is set by the spread cap, with the maker bottoming out well above the nominal floor in every calibrated cell.

Equally important is what this calibration does *not* claim. Spread and depth levels are not fitted: the stylized MM quotes a ~ 1.5 bp median spread where the perp trades at ~ 0.016 bp, and its depth-to-median-trade ratio is $5\times$ against the perp’s $\sim 2060\times$. The size distribution is not fitted in the experiments of §4: sim order sizes there are constants while real sizes are heavy-tailed (p99/p50 $\approx 670\times$). A lognormal size model for the calibrated re-run is implemented and parameterized ($\sigma_{\ln} = 2.80$ fitted on the p50/p99 ratio; median 1 unit with a calibrated book depth of ~ 2061 units, matching the real depth-to-median-trade ratio), with the disclosed compromise that a two-quantile lognormal under-predicts the p90 size ratio, $36\times$ against the

observed 59 $\times$. And the latency/race structure is not estimable from any public crypto feed — measuring races requires failed-order message data of the kind Aquilina, Budish, and O’Neill (2022) obtained from a regulator — so HFT latency and the race remain swept design parameters throughout. The calibrated claim is therefore deliberately narrow: the *environment* the agents trade in (arrival intensity, fundamental volatility) is disciplined by data; the strategic structure is designed, not estimated.

3.7 Measuring the snipe-to-depth ratio on the tape

§4.1 sweeps the snipe-to-quote size ratio, and where the real market sits on that axis is what turns the surface from a parameter study into a claim. No public feed reports the quantity, so we measure it.

The sample is 86 complete book-and-trade hours of BTCUSDT-perp — hours whose book file carries at least 30,000 snapshot rows against a nominal 36,000 at the collector’s 10 Hz cadence, and for which a matching trade file exists. Completeness is the binding requirement because every trade must be divided by the top-of-book depth it actually met, read from the snapshot immediately preceding it rather than the one following, which already reflects the trade’s own consumption. That yields 2.27 million taker orders, after regrouping Binance’s per-price-level aggregated-trade records back into single taker orders — a correction without which multi-level sweeps, precisely the race flow of interest, fragment and bias the ratio downward.

Nothing in a public tape flags a snipe, so we do not label them: we report the distribution unconditionally and under a *race proxy* — an order arriving within 0.2 s of a snapshot on which the top of book moved by at least one price increment — and treat the gap between the two as the informative object. Conditioning on a price move rather than on order size is deliberate, since a move is what the BCS mechanism says triggers a race whereas conditioning on size would assume the conclusion. A placebo on non-move orders in the same window quantifies the dilution that 100 ms snapshots impose on a microsecond phenomenon: its median is 0.0057 against the proxy’s 0.0283, a factor of 5.0, which is what licenses calling the proxy a bound rather than a measurement.

sample	n	p10	p50	p90	p99
unconditional	2,269,256	0.00027	0.0072	0.393	18.0
race proxy	390,540	0.00051	0.0283	4.24	290
placebo (no move)	1,849,504	0.00024	0.0057	0.208	2.18

The distribution is extremely right-skewed, which is why the surface sweeps as far as ratio 1.0. On the collector’s current 105-hour sample the unconditional quartiles are p25 0.0011 and p75 0.0493 around a median of 0.0068, with the race-proxy median stable to 0.3%, so nothing in §4.1 turns on the sample boundary. Full alignment rules, the sweep de-fragmentation and aggressor-attribution corrections, and the placebo construction are in the replication materials (Appendix B).

4. Results

We report the incidence surface first (§4.1) and its extension to competing makers and elastic demand (§4.2), then the four supporting experiments on which the incidence claim rests (§4.3) and their re-run at calibrated parameters, which supplies magnitudes (§4.4). Throughout, the zero-sum conservation residual closes to floating-point noise on every run, which shows the harness’s accounting has no leak — and no more than that (§3.4). A placebo-validated null from a Hawkes endogeneity diagnostic is in Appendix C.

4.1 The incidence surface: who bears the latency tax

Incidence is not pinned by the data. The calibrated main arm sets the snipe size equal to the quote size, preserving the operating point’s ratio of 1.0 so a snipe clears the top quote. Holding the snipe at 50 units instead — 2.4% of the book — changes not the size of the effect but its direction: total rent rises in k rather than staying flat, per-HFT rent flattens instead of decaying as $1/k$, and the maker’s PnL delta is *positive* for $k \leq 5$, turning negative only at $k \geq 8$, with noise traders bearing the loss throughout. The claim that rent is extracted from the market maker is therefore conditional on the snipe-to-quote size ratio. Two ratio points do not locate a boundary in a two-dimensional space, which is what this section maps. The direction of the effect — the maker profiting from an arms race conducted against it — is surprising enough to deserve a mapped region rather than a sensitivity note. We sweep snipe size against HFT count on the calibrated environment of §3.6, seven snipe quantities $\times$ nine HFT counts at $n=100$ seeds (6,300 treatment runs plus a shared 100-seed no-HFT baseline: 6,400 engine runs), holding the maker count at one.

Validity: the surface nests the published cells rather than re-specifying the environment. Because a new grid could differ from the published calibrated runs for uninteresting reasons, the ratio axis is swept as integer snipe *quantities* so that two cells coincide exactly with results already reported. The ratio-1.0 column reproduces the calibrated Experiment 3 run to every quoted digit — rent 85,721 / 85,163 / 84,375 / 83,262 / 81,469 / 78,696 / 74,261 and liquidity gaps 32.05 / 42.19 / 52.73 / 74.19 / 106.51 / 151.08 / 217.78 across $k = 1 \dots 21$ — and the ratio-0.0243 row reproduces the 50-unit-snipe arm above to the digit, including the maker’s +1,327 delta at $k=1$ and −5,213 at $k=21$. The zero-sum residual is machine zero in all 63 cells (largest 5.6×10^{-10} against rents of order $10^{4-10}5$), so the accounting closes here as it does in every other result.

The boundary. Incidence is classified per cell by the sign of the maker’s PnL delta against the shared no-HFT baseline, at 95% bootstrap confidence; a cell whose interval brackets zero is reported as indeterminate rather than assigned to whichever side its point estimate falls on.

Snipe/quote ratio	Maker gains up to	Flips at	Maker delta at $k=21$
0.0049	all $k \leq 21$	—	+2,154 (still gains)
0.0097	all $k \leq 21$	—	+1,425 (still gains)
0.0243	$k = 6$	$k = 7$	−5,213
0.0485	$k = 2$	$k = 3$	−17,203
0.1213	$k = 1$	$k = 2$	−39,006
0.3396	none	below $k=1$	−61,978
1.0000	none	below $k=1$	−72,332

At the race proxy reading of 0.0283, the flip occurs at $k=7$. If the real market operates with more than seven competing snipers at that ratio, the inversion does not hold under the race proxy.

The critical count falls monotonically as the snipe grows, from beyond the grid at the smallest ratios to below one at the largest — a collapse spanning the full range of k over roughly one order of magnitude in ratio. At the two smallest ratios no flip occurs anywhere on the grid: the maker’s gain does not merely survive competition, it *increases* in k (from +1,061 at $k=1$ to +2,154 at $k=21$ at ratio 0.0049). At the two largest the maker loses even to a single sniper. The two calibrated snipe sizes above were one slice each of this surface.

Where the real tape sits, and what it does to the BCS incidence prediction. The ratio is a design parameter in the sense that no feed reports it directly, but §3 bounds it:

measured per trade against contemporaneous top-of-book depth on 86 complete book-and-trade hours of BTCUSDT-perp, it is 0.0072 at the median unconditionally and 0.0283 under the race proxy, with the placebo quantifying the dilution. The unconditional median falls between the two smallest ratios on the grid — and that is the region in which the maker gains at every HFT count we tested, while noise traders bear the entire transfer. **In the regime the tape supports, under single-maker liquidity supply, the BCS prediction that the latency tax falls on liquidity providers inverts.** That qualifier is load-bearing and is discharged in §4.2, which adds the maker-count axis and finds the inversion does not survive a second maker; this subsection’s conclusion should not be read on its own. The maker widens in response to adverse fills, and that widening leaves it better off on net — by cutting what the snipers take from its stale quotes rather than by earning more from noise flow, which in fact falls (see the leg decomposition below); the loss lands on the noise traders who cross the wider spread. The race-proxy median of 0.0283 sits just above the 0.0243 row, where the flip occurs at $k=7$ — so under the more conservative conditioning the maker still gains against a handful of snipers and loses only once the field grows. Both readings place real BTC at or below the boundary rather than deep in the maker-pays regime that §4.3 reports, and the maker-pays regime is reached only at snipe sizes 35–140 $\times$ the measured median. One qualification attaches to the measurement: the tape under-measures multi-level sweeps, because Binance aggregates fills per taker order per price level, so the true race-order ratio is somewhere above the proxy — which is precisely why the boundary matters more than any single point on it.




Figure 1: Incidence boundary over snipe-to-quote ratio and HFT count. The shaded vertical band is the ratio measured on the BTCUSDT-perp tape (unconditional median to race-proxy median); open markers with arrows are ratios at which no flip occurs anywhere on the swept grid.

The boundary is plotted with its censoring made explicit, because only three of the seven ratios flip inside the grid: at 0.0049 and 0.0097 the maker gains at every count up to twenty-one, so

the critical count is bounded below by 21 rather than measured, and at 0.3396 and 1.0000 the maker already loses to a single sniper, so it is bounded above by 1. Drawing a curve through all seven would fabricate four of them. The figure’s substantive content is the position of the tape band relative to the boundary: the band lies entirely within the maker-gains region across the full range of counts we swept, and for the BCS incidence prediction to hold on this market the measured ratio would have to sit above the boundary curve rather than below it.

A secondary regularity: the noise-trader toll is nearly invariant. Across all 63 cells the noise traders’ delta stays within $-1,699$ to $-3,410$, a factor of two, while the maker’s delta swings from $+2,154$ to $-84,023$, a factor of forty in magnitude and a change of sign. The maker’s exposure is therefore essentially the whole source of variation on this surface. Noise traders pay a roughly fixed toll for transacting against a latency-disadvantaged maker in the presence of snipers, and what the ratio and the count determine is whether the maker adds to that toll or collects part of it. This is a cleaner statement of the welfare picture than a single-cell decomposition can give, and it survives the conservation check in every cell.

The near-invariance is what inelastic demand predicts. Noise-trader arrival in this section does not respond to the spread: each agent crosses with a fixed probability per tick whatever the maker is quoting (§3.2), so a near-constant toll across 63 cells is the signature that assumption leaves. Volume is pinned by construction, so the only channel through which widening can reach the noise traders is the price paid per crossing, and the factor-of-two range is that channel operating alone. We therefore cannot separate the reading in which the toll is invariant for economic reasons from the reading in which the model gives it no margin on which to vary. §4.2 supplies the missing margin by making arrival elastic in the quoted half-spread, and finds the incidence unmoved.

A negative result: the surface does not reduce to one dimension. The boundary invites an obvious simplification — that only aggregate snipe pressure matters, so incidence should turn on the product of ratio and count rather than on each separately. At the flips the product ratio $\times k$ lies in a narrow band, roughly 0.10–0.20, which makes the reduction look right. We tested it and it fails. At ratio 0.0097 with $k=21$ the product is 0.204 and the maker still gains significantly; at ratio 0.0243 with $k=7$ the product is 0.170 — smaller — and the maker no longer gains. A single product threshold cannot order those two cells, so the near-collapse is a coincidence of the grid rather than a law, and the structure is genuinely two-dimensional. We report no product rule. The mechanism is presumably that snipe size and sniper count act through different channels — size determines how much of the quote a single race removes, count determines how often the quote is stale when a race starts — and the maker’s widening response is not linear in either.

Decomposing the inversion into its two legs. The account given above — that widening leaves the maker better off by cutting what snipers take from its stale quotes while its noise-flow revenue falls — is a claim about two quantities, and the surface reports only their difference. We therefore measured the legs separately at the tape-band cells ($n = 100$, single maker as in the rest of this section). Every maker fill is scored as edge against the fundamental at fill time, positive when the maker captured value, and partitioned by the *same* adverse test that triggers the widening (§3.2), so the decomposition is measured with the mechanism’s own discriminator rather than an external proxy. Because maker PnL is mark-to-market while the legs are realized edge at fill time, inventory carry is reported as an explicit third column rather than absorbed into either leg; the three sum to the net delta exactly, and the net column reproduces the surface’s ΔMM (largest reconciliation error 3.6×10^{-12}).

ratio	k	revenue leg	sniping leg	inventory carry	net Δ MM
0.0049	1	-2,949	+2,289	+1,721	+1,061
0.0049	3	-5,481	+4,787	+2,282	+1,588
0.0049	7	-6,487	+6,040	+2,360	+1,913
0.0049	21	-6,747	+6,526	+2,374	+2,154
0.0097	1	-4,329	+3,512	+2,062	+1,244
0.0097	3	-6,249	+5,448	+2,343	+1,542
0.0097	7	-6,852	+6,041	+2,369	+1,558
0.0097	21	-6,529	+5,566	+2,389	+1,425

Seed-matched deltas against the shared no-HFT baseline, $n = 100$.

The revenue leg is **negative** in all eight cells: with snipers present the maker earns *less* from non-adverse flow than it does in the no-HFT baseline, because widening drives away more noise volume than it recovers in edge per fill. The sniping leg is **positive**, meaning the maker’s losses to adverse fills *shrink* relative to baseline — the wider quote is picked off less profitably, and the baseline itself carries adverse selection from noise traders hitting stale quotes when the fundamental moves. The net gain is therefore not “more noise revenue than sniping loss.” It is reduced adverse-selection cost, plus a substantial inventory-carry term, together outweighing a fall in noise revenue.

Two consequences follow. The inversion is real and is not an artifact — the net column is the surface’s own Δ MM and it is positive throughout — and the intuitive story in which widening earns the maker more from noise flow is the wrong way round on the revenue leg. And the inventory-carry column is large enough (roughly a third to a half of the gross legs’ difference) that a decomposition omitting it would misattribute the source of the gain; the maker’s benefit is partly that it holds inventory through a favourable mark, which is a distinct channel from either leg and one the widening story does not describe at all. Which of the three channels dominates under elastic demand is not yet measured.

The claim holds for all HFT counts tested — up to twenty-one. Behavior at higher counts is not measured.

4.2 Structural robustness: maker competition and elastic demand

§4.1 holds two things fixed that its own limits identify as load-bearing: liquidity is supplied by a single maker, and uninformed demand is inelastic. This section relaxes each in turn. The maker-count axis reverses the inversion at the class level; elasticity does not move it. Both extensions nest §4.1 exactly at their inelastic, single-maker corner, checked numerically rather than asserted.

Maker competition. We re-run the 63-cell surface at maker counts $M \in \{1, 2, 3, 5\}$ on the calibrated environment, holding every §3.6 parameter fixed, at $n = 100$ (441 cells, 44,800 engine runs). Aggregate top-of-book depth is held at the calibrated quote size and divided across the makers, so adding makers varies competition and latency spread rather than the depth a snipe meets; latencies for $M > 1$ are §4.4’s set, anchored on measured replenishment quantiles. A zero-latency-leader variant, placing the leading maker at zero latency, is run as an optimistic bound. The $M = 1$ column reproduces §4.1 to **0.0 absolute difference** on five metrics at three probe cells re-run at $n = 100$, and conservation holds across all 441 cells at a maximum absolute residual of 5.57×10^{-10} .

M	latencies (ticks)	ratio 0.0049	ratio 0.0097	flip @ 0.0243	rent at $k = 3$, ratio 1.0
1	3	gains at all $k \leq 21$	gains at all $k \leq 21$	$k = 7$	84,375
2	1, 5	gains at $k \leq 2$	gains only at $k = 1$	none	22,645
3	1, 3, 5	no gains at any k	no gains at any k	none	9,575
5	1–5	no gains at any k	no gains at any k	none	2,499

The inversion all but fails at $M = 2$, the smallest competitive case, and fails completely at $M = 3$. Of eighteen tape-band cells, eighteen are maker-gains at $M = 1$, three at $M = 2$ — only at the lowest HFT counts, $k \leq 2$ at ratio 0.0049 and $k = 1$ at 0.0097 — and none at $M = 3$ or $M = 5$; the zero-latency-leader arms show none at any $M > 1$. This is a sign change carried by the intervals rather than a loss of power: at $k = 3$, ratio 0.0097, the maker’s delta runs +1,542 [1,278, 1,844] at $M = 1$, −741 [−995, −477] at $M = 2$, −613 [−694, −533] at $M = 3$, −189 [−306, −54] at $M = 5$. The noise-trader toll shrinks in step, from −2,759 under a single maker to −205 at five, so competitive supply reinstates the incidence BCS predicts rather than merely cancelling the inversion. Two channels drive it and both are real: splitting aggregate depth multiplies each maker’s own snipe-to-quote ratio by M — the aggregation problem of §5.2 measured rather than argued — and, at matched per-maker ratio, competition still moves the boundary (at per-maker 0.0243 the single maker gains through $k = 6$ while $M = 5$ already pays from $k = 5$).

The aggregate conceals a within-class redistribution, and this is the finding that survives. A negative summed delta is consistent with every maker losing or with some gaining while others fund them. Re-running the tape-band cells with each maker’s delta bootstrapped separately, at $n = 100$, across eight latency configurations (maximum residual 5.57×10^{-10}), one maker gains while another pays in 136 of 144 cells, each verdict at 95% bootstrap confidence; in 4 cells every maker pays. Who pays depends on the configuration:

makers’ latencies (ticks)	fastest gains	middle gains	slowest pays
0, 5 (zero-latency leader)	18	—	18
0, 3, 5 (zero-latency leader)	18	18	18
1, 5 (calibrated)	14	—	15
2, 6	11	—	14
1, 4	15	—	17
1, 3, 5 (calibrated)	4	18	13
1, 4, 7	6	18	14
2, 3, 4	6	3	3

Cells out of 18 per row. With a zero-latency leader the ordering is monotone in latency: at $M = 2$ the leading maker gains +2,284 [2,096, 2,477] at $k = 3$, ratio 0.0097, while the five-tick maker pays −3,475 [−3,656, −3,303]. With latencies spaced as measured replenishment spaces them, it is not. At two makers the faster usually gains, but the order reverses at thirteen or more snipers. At three makers without a zero-latency leader the fastest never wins consistently:

the middle maker gains in all eighteen cells under $\{1, 3, 5\}$ and $\{1, 4, 7\}$, while under the closely spaced $\{2, 3, 4\}$ the *slowest* gains most often, in 10 of 18, and the middle maker pays in 14. Neither speed rank nor a particular tick value predicts the winner — three ticks wins under $\{1, 3, 5\}$ and mostly pays under $\{2, 3, 4\}$.

Agent order is ruled out as the cause. Re-running the calibrated sets with the latencies assigned to the maker slots in every order leaves every verdict unchanged: at three makers, where depth splits evenly, the per-maker results are identical to the bit across all six orders, and at two they differ by at most 276, through the one extra unit the first slot quotes. The identity orders reproduce the published per-maker run to 0.0 on all ninety per-maker means. §4.4’s monotone-in-latency ordering was measured at ratio 1.0 and does not extend to the tape band, and what sets the winner there we have not identified.

Elastic demand. The second fixed assumption is that noise traders cross at a fixed per-tick probability whatever the maker quotes. The objection this licensed — elasticity withdraws the flow that finances the maker’s gain — rests on a premise §4.1’s leg decomposition already refuted: the revenue leg is *already negative*, and the gain comes from the sniping leg shrinking. Elasticity therefore deepens a leg already working against the maker without touching the dominant channel, making the question quantitative. We set $\lambda_{\text{eff}} = \lambda(s_{\text{base}}/s)^\alpha$ with s read off the book, so arrivals match §3.6’s fitted intensity at base spread for every α and only the response to widening changes, and re-run the full surface at $\alpha \in \{0, 0.5, 1.0\}$ (189 cells). The $\alpha = 0$ path short-circuits before the book is read, preserving the RNG stream: it reproduces §4.1 to 0.0 absolute difference on rent and the zero-sum residual at all three probes, with maximum residual 1.14×10^{-8} and baseline trades falling $5,165 \rightarrow 4,902 \rightarrow 4,657$, confirming the knob engages.

α	ratio 0.0049	ratio 0.0097	flip @ 0.0243	rent at $k = 3$, ratio 1.0	baseline trades
0.0	gains at all	gains at all	$k = 6$	85,638	5,165
(inelastic)	$k \leq 21$	$k \leq 21$			
0.5 (mild)	gains at all	gains to	$k = 5$	86,931	4,902
	$k \leq 21$	$k = 13$			
1.0 (unit)	gains at all	gains to	$k = 5$	88,502	4,657
	$k \leq 21$	$k = 13$			

This surface is at $n = 20$, so its $\alpha = 0$ row differs slightly from the $n = 100$ single-maker figures above (flip at $k = 7$, rent 84,375); comparisons across α are within the one run. The maker-gains region survives. Eighteen of eighteen tape-band cells gain inelastically and **seventeen at both mild and unit elasticity**; the single loss, ratio 0.0097 at $k = 21$, turns *indeterminate* rather than maker-pays (+215 [−202, 672] at $\alpha = 1.0$), and no band cell classifies as maker-pays at any elasticity tested. The lower band edge is untouched, the 0.0243 boundary retreats one grid step, and the effect saturates between $\alpha = 0.5$ and $\alpha = 1.0$ rather than scaling. This is an out-of-sample test of the mechanism correction rather than a restatement of it: on a revenue-capture account unit elasticity should have been close to fatal, since it removes a tenth of noise participation exactly when the maker widens. Noise traders do withdraw as the objection predicted — their loss shrinks from −2,667 to −2,033 — which buys them partial self-protection but does not move the incidence.

What the two extensions jointly establish. The class-level answer is structure-dependent, and so is the within-class one. Whether summed maker PnL is positive depends on how many makers supply the book; that the class splits into winners and losers holds in 136 of 144 cells, but who is on which side depends on how the makers’ latencies are spaced. Elasticity was tested at

a single maker, so it supports the single-maker inversion and says nothing yet about the within-class split. The maker-count surface is mapped at $n = 100$, the standard of the within-class decomposition; the elasticity surface remains at $n = 20$.

4.3 What the incidence claim rests on: Experiments 1–4

Four experiments at the pre-registered operating point establish the properties the incidence surface assumes: that the transfer is real and its accounting closes, that the mechanism carries no hidden threshold, how competition among snipers changes the rent pool and the book, and that the clearing rule drives the outcome. Per-cell tables are in Appendix D.

Experiment 1: welfare decomposition and the zero-sum conservation check We begin by checking that the welfare accounting closes. By construction, every unit of profit-and-loss in the model-checked matching engine accrues to one of three agent classes — high-frequency traders (HFT), market makers (MM), or noise traders (NT) — so the sum of realised PnL across classes must equal zero on every run. We measure the empirical conservation residual across 100 seeds at the operating point ($n_{\text{hft}} = 3$, adverse sensitivity 0.1, spread decay 0.1, MM latency $3000 \mu\text{s}$). In the treatment economy the residual closes to machine zero, $2.78\text{e-}12$ $[-3.34\text{e-}12, 9.11\text{e-}12]$, and the matched no-HFT baseline behaves identically, $-5.39\text{e-}12$ $[-2.16\text{e-}11, 1.05\text{e-}11]$. Because the residual is indistinguishable from floating-point noise in both arms, the decomposition of welfare into HFT rent, MM loss, and NT loss is not subject to unmodelled leakage. That is a check on bookkeeping only: every fill is bilateral and every agent is marked at one price, so any correct accounting closes, and the residual says nothing about whether the matching or the economics is right.

Within this conserved frame, HFT entry extracts a positive rent of 34.70 $[26.25, 44.80]$ per run. The transfer is borne almost entirely by market makers: MM PnL changes by -28.50 $[-38.31, -20.35]$ relative to the matched baseline (from 101.27 $[93.28, 110.72]$ to 72.78 $[71.05, 74.38]$), while the noise-trader class absorbs a much smaller share, with an NT PnL delta of -6.21 $[-7.87, -4.61]$. At $n = 100$ this NT delta is now statistically distinguishable from zero — at $n = 20$ its interval straddled zero — but it remains an order of magnitude smaller than the MM delta and accounts for roughly 18% of the transfer. The welfare transfer is therefore predominantly, though not exclusively, borne by liquidity suppliers, consistent with HFT rent arising chiefly from picking off stale market-maker quotes rather than from worsened execution for uninformed flow.

The microstructure correlates of this transfer are pronounced. Liquidity gaps more than triple under HFT entry, 9.70 $[9.15, 10.27]$ in treatment against 2.84 $[2.47, 3.26]$ in the matched no-HFT baseline, a delta of 6.86 $[6.20, 7.52]$. Lagged Kyle’s λ , our measure of price impact per unit signed flow, rises by roughly an order of magnitude, from 0.200 $[0.184, 0.216]$ to 1.372 $[1.361, 1.383]$; the coefficient takes its name and its interpretation from Kyle (1985), where λ is the inverse of market depth in a continuous-auction model with a strategic informed trader, so a rise in it is properly read as the book becoming easier to move per unit of signed order flow. Table D1 collects these estimates.

We stress-test the rent result against the two parameters that most directly govern the adverse-selection channel, sweeping adverse sensitivity and spread decay over the grid $\{0.05, 0.1, 0.2\} \times \{0.05, 0.1, 0.2\}$ with $n = 100$ seeds per cell. HFT rent is significantly greater than zero in all nine cells, and the zero-sum identity holds in all nine, with point estimates of the rent ranging from 16.56 to 46.63 . The rent is thus a robust feature of the calibration rather than an artefact of a single operating point.

Two points of nuance bear on how the result should be read. First, the no-HFT baseline is mildly fragile: liquidity gaps in that arm are not exactly zero (2.84 $[2.47, 3.26]$), reflecting residual quote churn even absent HFT participation, so the baseline does not represent a frictionless market.

We nonetheless interpret the treatment-versus-baseline contrast as informative because the gap CIs are non-overlapping (treatment 9.70 [9.15, 10.27] versus baseline 2.84 [2.47, 3.26]). Second, the NT PnL delta is small but, at $n = 100$, statistically distinguishable from zero (-6.21 [-7.87, -4.61]); we therefore state that HFT entry imposes a measurable but secondary cost on noise traders, roughly 18% of the total transfer against the market maker’s 82%. The welfare story of Experiment 1 remains a predominantly market-maker-to-HFT transfer occurring under a conserved, machine-zero accounting identity, accompanied by a sharp deterioration in measured liquidity.

Experiment 2: latency-advantage scaling Experiment 2 sweeps the market maker’s staleness — the latency with which it observes the fundamental — holding the HFTs at zero latency, and asks how the transfer scales. The question matters for the incidence surface: a boundary mapped across a mechanism containing a discontinuity would be an artifact of the discontinuity rather than a boundary.




Figure 2: HFT rent versus MM staleness with linear-fit overlay and significance threshold

Rent rises with staleness at approximately 10.46 per tick, and the relationship is smooth: comparing a linear specification against a segmented one with a free breakpoint gives $\Delta\text{BIC} = -0.17$, a magnitude well inside the range of indifference, so the data do not discriminate between them. We report this as the absence of a detectable knee rather than as positive support for linearity — a distinction that matters, because at our sample size an absence of evidence for a threshold is not evidence of its absence. Rent first becomes statistically distinguishable from zero at roughly two ticks of staleness, which is the detectability floor of this grid rather than a property of the mechanism, since latency resolves only at tick granularity here (§5.2).

The welfare decomposition closes at every point on the sweep, and liquidity gaps and HFT fill rate move with staleness in the directions the mechanism predicts: a staler maker is picked off more often, and its thinned book produces more one-sided episodes. Per-cell values across the sweep are in the replication materials (Appendix B); the figure carries the shape, which is what the incidence argument uses.

Experiment 3: HFT competition and the fragility of the book We next ask how the number of competing high-frequency traders, k , shapes the rents they extract and the costs they impose. We sweep k over the grid $\{0, 1, 2, 3, 5, 8, 13, 21\}$ with $n = 100$ seeds per cell, holding the market-maker (MM) population and order-flow process fixed. Two results come out of it, and they have different standing. That total rent stays flat as k grows is a consistency check rather than a finding: at this operating point one snipe takes the whole quote (snipe size equal to quote size), every HFT has the same latency and one wins each race on a uniform jitter draw, so the number of snipe opportunities is set by the maker’s staleness and the fundamental path and k can only divide it. That the book grows more fragile as k grows is a finding, and one the maker’s depth-thinning rule does not account for.

Per-HFT rent falls monotonically and steeply, from 34.82 [26.40, 44.84] at $k = 1$ to 1.60 [1.21, 2.07] at $k = 21$ (Table D2). The decline is steep and close to hyperbolic, and the same structure appears in execution shares: the per-HFT fill rate falls from 0.2886 [0.2863, 0.2909] at $k = 1$ to 0.0139 [0.0137, 0.0140] at $k = 21$. The fill rate times k is 0.289 at every count, to within 1% — what one winner per race among k symmetric entrants produces by construction.

The aggregate confirms this directly. Total extracted rent is flat across the entire grid: 34.82 [26.40, 44.84] at $k = 1$, 34.93 [26.55, 44.87] at $k = 2$, 34.70 [26.25, 44.80] at $k = 3$, 33.83 [25.36, 44.03] at $k = 5$, 34.50 [25.84, 44.76] at $k = 8$, 34.43 [25.63, 44.87] at $k = 13$, and 33.70 [25.39, 43.45] at $k = 21$ — a derived mean of 34.42 (std 0.48, range 33.70–34.93). MM PnL is approximately constant near 73 for all $k \geq 1$, and the welfare delta tracks total rent and is flat. The pool is fixed only because a single snipe clears the quote. When it does not — the arm of §4.1 with snipes of 2.4% of calibrated depth — later HFTs fill what the first left, and from $k = 1$ to $k = 21$ total HFT fills rise elevenfold and total rent sevenfold (1,179 to 8,623). Competition enlarges the pool once snipes are smaller than the quote; the full-quote arm here is the special case, and §4.1 maps where between the two the real market sits.

Where the marginal cost of entry does appear is in market fragility. The incidence of liquidity gaps rises steadily with k , from a baseline of 2.84 [2.47, 3.26] at $k = 0$ to 47.47 [46.22, 48.85] at $k = 21$ — roughly a 16.7-fold increase. The growth is well described as linear in k (derived fit: slope = 2.13, intercept = 3.74, $R^2 = 0.99730$). The maker’s thinning rule — its resting size shrinks as its spread widens — is the obvious mechanical explanation: more adverse fills, more widening, a thinner and more often one-sided book. It is not the explanation. Re-running the sweep with the rule off and quote size held constant ($n = 100$) leaves the growth in k unchanged: gaps run from 3.61 at $k = 1$ to 47.18 at $k = 21$, a fitted slope of 2.18 against 2.13 ($R^2 = 0.998$). Thinning adds about one and a half gaps at small k and nothing to the slope. Nor is the growth a count of snipes: fills stay flat (252 to 242 at fixed depth) while gaps rise thirteenfold, the book is sampled once per tick, and losing HFT orders are IOC and leave nothing behind. We have not identified the channel, and report the growth as a regularity with an unexplained mechanism rather than as emergent fragility.

No cost of speed is modelled. BCS call the arms race socially wasteful because firms spend real resources on speed. Nothing in this model spends anything: latency is a parameter, not an investment, so the transfer is zero-sum by construction and the race’s waste is outside what these experiments can measure. What Experiment 3 speaks to is how the transfer and the book’s fragility change with the number of fast traders.

Experiment 4: batch auctions as a remedy Experiment 4 switches the clearing rule from continuous matching to a uniform-price call auction, holding every agent and market parameter fixed, and sweeps the clearing interval. Both grids clear through the verified engine’s own uncross (Appendix A), so the comparison isolates the clearing rule rather than a reimplementations of it.

Market-maker welfare recovers monotonically once the interval is long enough for the slow quoter




Figure 3: Welfare decomposition versus clearing interval

to participate: at-least-50% recovery by an interval of about 25 ticks, and about 97% by 500, taking maker loss from 28.5 to 0.7 — the same under both tie-break rules the experiment has been run with (Appendix A). Over 50,000-tick runs the size of the recovery depends on the rule. Under the engine’s, every interval ≥ 5 recovers effectively all of the loss and intervals 10 and 25 leave the maker significantly ahead; under the retired rule the maker recovers 40–68% at intervals 5–25 and still loses. Both rules agree on the direction: batched, the maker does better than continuous at every interval of five ticks or more. That comparison, resting on a matched-seed decomposition rather than on the rent estimator, is the robust signal in the experiment.

The accompanying reduction in HFT rent points the same way but is not established. At 3,000 ticks only the interval-5 cell’s rent interval excludes zero, and at 50,000 ticks the estimator’s dispersion — which grows super-linearly in run length — swamps every batch cell above interval 1, and under the retired tie-break rule three of those cells were significant, so we draw no conclusion from the long grid and decline to nominate an optimal interval on rent grounds. At intervals above 100 ticks those cells are uninformative and are in the replication materials (Appendix B).

The interval-1 cell runs the other way: a per-tick auction worsens maker loss to 41.3, against 28.5 under continuous clearing, and liquidity gaps spike to about 137 against a continuous baseline of 9.7. The obvious reading — the maker’s three-tick latency exceeds the one-tick interval, so it cannot refresh within a clearing cadence — is close to a tautology, and it is not what the data show ($n = 100$, each arm against its own continuous control). A seven-tick maker facing a five-tick auction, slower than the interval, still does 40% better batched than continuous (−40.1 against −66.1). A maker exactly as fast as the interval (one tick, one tick) is worse under the per-tick auction by point estimate (−16.8 against −10.3), but the intervals overlap; and a zero-latency maker removes the race altogether, the rent being near zero in both arms. The backfire is specific to the per-tick auction rather than to latency exceeding the interval, and whether it survives a maker as fast as the interval is not established. Batching is a remedy above a

minimum cadence; what sets that cadence we cannot yet say.

4.4 Experiments 1–4 at calibrated parameters

The pre-registered operating point was chosen for mechanism clarity, not realism. §3.6 fits arrival intensity and fundamental volatility to BTCUSDT-perp moments, and this section re-runs Experiments 1–4 at those fitted values with a lognormal size model and the perp’s book depth. Every qualitative finding survives the move: the transfer remains a machine-precision zero-sum, rent scales smoothly in maker staleness with no phase transition, fragility rises in HFT count, and batch auctions recover maker welfare above a minimum cadence while the per-tick auction backfires. That survival is the strongest evidence we can offer that the findings are properties of the clearing design rather than of a convenient parameterisation.

What calibration adds is magnitude. Normalised by traded notional, HFT rent moves from 0.125 bp at the operating point to 4.42 bp calibrated. The calibrated figure is an upper bound rather than an estimate, because a single latency-disadvantaged maker facing the perp’s order flow has no competitor to replenish a cleared quote and so absorbs the entire race. Replacing it with N makers whose latencies span the measured top-of-book replenishment quantiles, holding aggregate depth fixed, collapses the rent by a factor of fifty:

N makers	latencies (ticks)	HFT rent (bp)	liquidity gaps
1	3	4.421	52.73
2	1, 5	0.801	38.00
3	1, 3, 5	0.324	16.29
5	1–5	0.089	2.71

$n = 100$ seeds; the $N = 1$ row reproduces the published single-maker figure, so the sweep nests it rather than re-specifying it, and the zero-sum residual stays at machine zero throughout. The ≈ 0.4 bp latency-arbitrage tax that Aquilina, Budish, and O’Neill (2022) measure from regulator-held exchange messages falls inside the range these arms span — envelope consistency between a simulated and a directly measured quantity, not agreement on a number.

The mechanism is visible in the per-maker decomposition. At $N = 5$ maker PnL is monotone in latency, the fastest earning +2,014 and the slowest losing −2,393, because a fast competitor replenishes the touch before the slow maker’s stale quote can be sniped. The slow maker still absorbs adverse selection; what changes is that it is no longer the only supplier. The single-maker absorption bias is therefore measured and removed rather than asserted as a caveat.

At calibrated parameters the batch arm also remains non-monotone in the clearing interval, the interval-1 cell still the worst performer on both rent elimination and gaps. A second calibrated snipe size, 50 units, is the starting point of §4.1.

5. Discussion

The incidence results (§§4.1–4.2) rest on supporting experiments that establish a zero-sum transfer, smooth scaling in the maker’s staleness, fragility that rises with the number of snipers, and a clearing rule that drives the outcome (§4.3). Below we draw out what the findings do and do not license, state their limitations plainly, and set out the work needed to harden them.

5.1 Implications for market design

Three findings carry policy weight, and each is narrower than the headline it replaces.




Figure 4: Rent elimination against liquidity gaps across batch clearing intervals; the dotted line is the continuous arm’s gap rate

First, the class-level answer to “who bears the latency tax” is structure-dependent, so a regulator cannot read it off the existence of the tax. Whether summed maker PnL is positive or negative turns on how many makers supply the book and on the snipe-to-depth ratio prevailing there — quantities that differ across venues and that no public feed reports — so an incidence claim transported from one market to another is an assumption rather than a measurement.

Second, within-class redistribution survives the maker-count test, a permutation of agent order and a change of latency values, and is the more robust object. At the tape-measured ratio one maker gains while another pays in 136 of 144 cells; which one pays depends on how the makers’ latencies are spaced, and the split has not yet been tested against elastic demand. The policy reading is that latency competition sorts liquidity providers rather than taxing them uniformly: a remedy justified as protecting “liquidity providers” protects some of them at the expense of others, and which ones depends on their relative speeds.

Third, the clearest design warning is the interval-1 backfire. A per-tick auction does not merely fail to help — it raises maker loss to 41.3 against the continuous 28.5 and spikes liquidity gaps to about 137 against 9.7. It is not simply that the maker is slower than the interval: a seven-tick maker still gains from a five-tick auction (§4.3). Batching is a remedy only above a minimum cadence, and a venue setting the interval too aggressively can reproduce the pathology it meant to cure; what fixes that cadence is not yet measured. That the recovery signal is monotone above that floor while the accompanying rent reduction is not statistically resolved (§4.3) is itself the reason to frame the remedy in terms of who is protected rather than of the sniper’s take.

5.2 Limitations

Several limitations bound the strength of these conclusions.

First, the elasticity of uninformed demand is the assumption the incidence result

leans on hardest, even now that it is measured. Noise traders cross at a fixed per-tick probability in every arm through the competition surface (§3.2), so the volume a maker earns its widened spread on cannot fall when it widens. §4.2 relaxes this to unit elasticity and the maker-gains region very largely survives: eighteen of eighteen tape-band cells gain inelastically and seventeen at both mild and unit elasticity, the single loss turning indeterminate rather than maker-pays, with the 0.0243 boundary retreating one grid step and the lower band edge untouched. That resolves the limit in the direction §4.1’s leg decomposition predicts — elasticity deepens an already-negative revenue leg without touching the loss-reduction channel that produces the gain — but it does not eliminate it. Elasticity above unit is untested, as is its interaction with the maker-count axis, where the within-class split has not been tested against elastic demand at all.

Second, liquidity supply is narrow in both count and rule. §4.2 maps maker counts to five, but §4.1’s headline arm holds one maker, and every maker in this paper follows a single imposed adverse-selection quoting rule. A maker that defended itself differently — fading quotes, skewing on inventory — might not profit from small snipes at all, and none of ours chooses its spread to maximise profit, so “the maker gains” is a property of that rule under sniping rather than a statement about profit-maximising liquidity provision. The rule’s aggressiveness is one hand-tuned parameter, and re-running §4.1’s load-bearing readings across a $7.5\times$ span of it leaves the inversion’s *existence* intact at the tape band while moving its *location* across the whole grid and changing total rent by a factor of five; no critical HFT count should therefore be read as a property of the market rather than of this maker. Scope is narrow in ways we note rather than enumerate separately: one asset on one venue with no cross-venue routing or multi-asset correlation; spread and depth *levels* far from the perp’s even though the flow is fitted (§3.6); latency resolving only at tick granularity, so the two-tick detectability floor is a property of that grid; and IOC-in-batch semantics inside the auction arm that are a modelling choice rather than an institutional given, which is part of why the batch-arm rent reduction stays suggestive rather than established.

Third, the ratio measured on the tape and the ratio swept in the simulation are not the same object. The tape’s 0.0072 is a snipe against *aggregate* top-of-book depth — every maker resting at that level — whereas the simulation measures it against *one* maker’s quote. The per-maker ratio is therefore bracketed below by the tape ratio and above by the tape ratio times the number of makers at the touch, a count no public feed reports, since depth arrives pre-aggregated with no maker identifiers. This matters because the inversion is a statement about the per-maker ratio: if the true value sits above the boundary at the relevant HFT count, the maker bears the tax after all. §4.2 resolves the bracket in the direction it warned of — splitting aggregate depth across M makers multiplies each maker’s own ratio by M , and the maker-gains region at the tape band is all but gone at $M = 2$ (three of eighteen cells, at the lowest HFT counts) and gone at $M = 3$ — so what remains unobservable is M itself, and the question is now how many makers rest at the touch rather than which quantity the simulation sweeps. One methodological point belongs here rather than in a footnote: we apply no deflated-Sharpe-style multiple-testing correction at the cell level, because the surface reports every cell of a grid fixed before the runs rather than the survivors of a search, and a complete map selects nothing.

Fourth, the model’s clock and its costs are coarse exactly where the rents are small. One tick is 100 ms (§3.6), a convention the calibration data forces: the book snapshots from which maker replenishment latencies are measured arrive every 100 ms. But races on real venues are decided in microseconds (Aquilina, Budish, and O’Neill, 2022) and crypto market makers react in milliseconds, so the model cannot represent the race its flow is calibrated against at the race’s own resolution. Arrival intensity and volatility could be refit at 10 ms from the millisecond trade timestamps; maker latency at that resolution cannot be measured from this collector and would have to be assumed. Fees are absent altogether. Binance’s standard schedule for USDT-

margined futures charges about 2 bp to makers and 5 bp to takers, lower at higher tiers and with maker rebates under its market-maker programmes — the same order as, or larger than, the 0.089–4.42 bp rents of §4.4. A fee schedule enters every leg of the decomposition: it taxes the snipers’ IOCs and the noise traders’ crossings, and charges or pays the maker per fill depending on tier, so it could move both the size of the rent and which participant nets positive. We have not modelled it, and every incidence result here is gross of fees.

What survives all four is the structural point: the identity of the class bearing the latency tax is not settled by the existence of the tax. It is a separate empirical question with a structural answer, and at realistic snipe sizes the answer is not the one the theory predicts.

5.3 Future work

The limitations above map directly onto an agenda for hardening these results.

Competing makers with heterogeneous latencies, and spread-elastic uninformed demand, were both expected to threaten the incidence result, and neither did in the way anticipated (§4.3): competition removed the class-level inversion while leaving the within-class finding standing, and elasticity cost the maker-gains region one of eighteen tape-band cells. What stays open on that axis is elasticity above unit and its interaction with the maker-count surface, together with re-running the elasticity surface at $n = 100$; both surfaces resolve their boundaries only to the width of the grid.

The first open priority is the real-data leg of the Hawkes diagnostic. The simulation side is complete (Appendix C): time-based windows, quote-update and trade processes, and a passing placebo, yielding a validated null at the operating point. Running the same pipeline on the BTC tape requires one new component: a depth-depletion definition of a “gap” suited to a book that never empties (e.g., top- k depth falling below a percentile floor), after which the windowing, labelling, and placebo machinery apply unchanged. That leg is where the endogeneity hypothesis gets its empirical test, in the Filimonov–Sornette spirit, on flow where self-excitation is known to exist.

Second, the single-venue, single-asset scope should be relaxed. A multi-venue extension would let us study latency arbitrage across correlated instruments and order routing, the setting in which the mechanical-arbitrage evidence motivating BCS is most salient, and would test whether the rising fragility of Experiment 3 survives when snipers can compete across books rather than within one.

Third, the channel through which liquidity gaps rise with HFT count (§4.3) remains unidentified, and the strategy spaces of both the market maker and the HFTs are deliberately narrow and should be enriched — for example, inventory-aware or quote-fade defences for the maker and more elaborate timing or order-type strategies for the snipers — to test how far the welfare decomposition and the batch-auction recovery are robust to more sophisticated play on both sides.

Fourth, the adverse-sensitivity check of §5.2 should be finished on its own terms, since it currently resolves the boundary’s direction more confidently than its location. Two refinements would do it: re-running it at §4.1’s $n = 100$ with the full ratio axis restored, the present $n = 20$ check pinning the boundary only to the width of the k -grid; and extending it to spread decay, which enters the same stationary-spread relation and was held at its pre-registered value throughout. The spacing point applies to §4.2’s surface too: at $n = 100$ it places the closing of the tape-band region between $M = 2$ and $M = 3$, a resolution set by the M grid.

6. Conclusion

The finding we regard as this paper’s most durable is about incidence *within* the liquidity-providing class, not about the class as a whole. At the snipe-to-depth ratio measured on the BTCUSDT-perpetual tape, competing makers split the incidence: across eight latency configurations at two and three makers, one maker gains while another pays in 136 of 144 per-maker cells (§4.2). Who pays is set by how the makers’ latencies are spaced rather than by speed rank — the slowest usually pays when latencies are far apart, and in the one closely spaced set we ran it wins most often — and the pattern is unchanged under every assignment of latencies to agents. The class-level answer is by contrast structure-dependent: a lone maker at the measured ratio profits from the arms race and noise traders bear the transfer, while adding a competitor reverses that at all but the lowest sniper counts and restores the incidence Budish, Cramton, and Shim predict (§4.2). Aggregate maker PnL is therefore a weighted sum over unlike participants rather than a verdict on liquidity provision, and the question “do liquidity providers bear the latency tax?” has no structure-free answer. The defensible statement is narrower and, we think, more useful: the tax is redistributed within the liquidity-providing class, and who bears it depends on the makers’ relative speeds. One condition on it deserves emphasis rather than a footnote: every maker here follows an imposed adverse-selection widening rule with an exogenous sensitivity parameter, so “the maker gains” is a property of that rule under sniping and not a claim about profit-maximizing liquidity providers who choose their spread endogenously — a maker free to re-optimize its quoting policy might capture more, or decline to quote at all, and nothing in this design bounds which.

That claim rests on a matching engine whose design is model-checked at bounded scope, on welfare accounting that closes in every cell, and on supporting experiments whose findings — and whose limits, from the unexplained fragility channel to the tie-break dependence of the batch result — are reported in §4.3.

The incidence result itself spans three structural axes. Sweeping snipe-to-quote ratio against HFT count under single-maker supply yields a monotone boundary rather than a single answer to “who pays”: at large snipe sizes the maker bears the transfer and Experiment 1 replicates (§4.3), while at small sizes it gains and noise traders bear the loss — through loss reduction rather than revenue capture, since the maker’s noise-flow revenue falls when it widens and the benefit comes from the accompanying drop in adverse-fill losses, with inventory carry supplying a third to a half of the net (§4.1). Adding the maker-count axis removes that inversion almost entirely when a second maker enters — it survives only at the lowest sniper counts — and completely at three, restoring the BCS prediction at the class level (§4.2); relaxing the inelastic-demand assumption up to unit elasticity leaves the single-maker picture intact, costing the eighteen tape-band cells only one, which turns indeterminate rather than reversing (§4.2). Measured on the perpetual’s own tape the ratio sits at 0.0072 at the median, and the calibrated rent envelope of 0.089–4.42 bp — from five competing makers to a single one (§4.4) — contains the ≈ 0.4 bp that Aquilina, Budish, and O’Neill measure from regulator-held exchange messages — envelope consistency between a simulated and a directly measured tax, not agreement on a number. The existence and sign of the rent were never in question here. What we add is that its incidence is a separate empirical question; that the answer depends on snipe size, sniper count, and how many makers supply the book; and that at the ratio the tape measures, the latency tax is redistributed within the liquidity-providing class rather than borne by it as a whole.

Appendix A. Verification and the differential test

A.1 Engine configuration

All trades in the continuous arm are produced by one C++ matching engine, `OrderMatcher::MatchingEngine`, the same object that is the subject of the project’s TLA+ specification suite. In the BCS

harness the engine is run in **synchronous** mode: `EngineHarness::start()` calls `engine_.start()` rather than `startAsync()`, so `submitOrder()` dispatches matching on the calling thread and the `OrderBook`'s `EventListener` fires inline. Draining trades immediately after a submit therefore returns exactly the fills that order produced, which keeps the harness deterministic and free of cross-thread or GIL concerns (`bcs_research/bindings/engine_bindings.cpp`, lines 5–16). The engine offers price-time priority matching, best bid/ask, depth snapshots, VWAP, and an optional circuit breaker; the breaker is disabled in this harness via `setCircuitBreakerThreshold(1e18)` when `start(disable_circuit_breaker=True)` so that simulation control, not an exchange-level halt, governs the dynamics (lines 81–83). Prices are fixed-point: a quoted price of 100.00 is `100 * PRICE_PRECISION`, with `PRICE_PRECISION` re-exported as `bcs_engine.PRICE_PRECISION` (line 186).

Critically, **the engine has no notion of agent latency**. The comment at `engine_bindings.cpp` lines 15–16 states this explicitly: the entire BCS latency/race mechanism lives in the Python scheduler above this layer. The engine's job is only to be a correct continuous double-auction matcher, and that correctness is what the spec suite certifies.

A.2 The specifications

The directly relevant specification is **Auction.tla** (`spec/Auction.tla`), which models the engine's session lifecycle over the seven legal `TradingStates` (`PreOpen`, `AuctionOpen`, `Continuous`, `AuctionClose`, `PostClose`, `Halted`, `VolatilityAuction`) and the discrete uncross that is the only way trades emerge from an auction state. Under `Auction.cfg` (`BROKEN_NO_REOPEN=FALSE`, `MaxOrders=3`, `MaxEvents=6`, `Prices={100,200}`), TLC checks the following **invariants**:

- **TypeOK** — the type invariant.
- **ValidState** — every reachable state is one of the seven legal `TradingStates`.
- **NoMatchOutsideContinuousExceptUncross** — no trade prints in an auction, halt, or closed state except a discrete uncross whose source is a true auction state (`PreOpen`, `AuctionClose`, or `VolatilityAuction`). Each trade carries an `inState` history field, which gives the invariant teeth.
- **SingleClearingPrice** — every uncross firing produces trades at exactly one price drawn from `Prices`. This is the code path the batch arm now clears through via the bridge's uncross (A.3).
- **NoTradeWhileHalted** — no trade record may carry `inState=Halted`.

As temporal/action **properties** it checks **ValidTransitions** (`[[] state'=state \/
<<state, state'>> \in LegalEdges]_state` — every step is either a self-loop or a legal directed edge of the session state machine) and the **liveness** property **VolatilityEventuallyReopens** (`[[] (state=VolatilityAuction => <>(state \in {Continuous,Halted}))`), proved under weak fairness `WF_vars(ReopeningUncross)`). Non-vacuity is established by a companion `AuctionBroken.cfg` (`BROKEN_NO_REOPEN=TRUE`) that disables the reopening uncross so TLC reports a liveness counterexample, confirming the property is not vacuously true.

We did not recover a model-checked state count for Auction.tla: a `states/` directory and TLC trace (`.bin`) artifacts exist in the spec tree, but we did not parse them, and no state count is reported for it. The large reachable-state count cited elsewhere for the engine (§3.1) pertains to `MatchingEngine.tla` and is recorded in `docs/Verification.md`, not derived from `Auction.tla`.

`Auction.tla` sits within a broader suite (`spec/`) that verifies the surrounding engine and infrastructure. We summarize the checked properties rather than restate every configuration:

Seven further specifications cover the matching core, replication, queue, and session lifecycle; their checked properties and configuration parameters are documented in the paper's replica-

tion archive. The one number we cite in the main body is for `MatchingEngine.tla`, whose deepest exhaustive configuration (`MaxOrders=4`, `MaxTime=2`, two prices) explores 171,187,419 distinct reachable states to completion with zero invariant violations, checking `FIFOExecution` and `MatchingConservation` for the continuous cross action alongside the book-management invariants. A companion broken configuration admits non-FIFO fills so TLC returns a counterexample, establishing non-vacuity.

The payoff is that any artifact the simulation reports — a fill, a price, a snapshot, a zero-sum residual — is produced by a matcher whose price-time priority, FIFO preservation, non-negativity, and single-clearing-price behaviour have been model-checked, not merely tested.

Two boundaries on that verification deserve equal precision, because “model-checked” is easy to over-read. TLC’s guarantee is exhaustive only within the stated bounds — a handful of orders, two prices, short horizons — so it certifies the matching *design* against its invariants at bounded scope rather than proving unbounded correctness. And the object verified is the TLA+ specification, not the C++ source: no refinement proof links the two, and the correspondence rests on construction and conformance testing. Wherever this paper says “verified engine,” both qualifications are attached. What the provenance buys, and what the internal-validity argument uses, is the elimination of design-level matching errors — the kind that could manufacture spurious gaps or transfers — within the checked scope, on the code path every trade in both arms traverses.

A.3 The bridge

A pybind11 module named `bcs_engine` bridges the verified C++ engine to Python (`engine_bindings.cpp`, `PYBIND11_MODULE` at line 212). It exposes a single entry class, `EngineHarness`, with: `add_symbol` (which must precede `start()`; `start()` freezes the book registry, after which `add_symbol` throws/no-ops, lines 64–70), `start(disable_circuit_breaker=True)`, `stop`, `submit_order(symbol, order_id, participant_id, side, price, quantity, order_type=Limit, tif=GTC)`, `submit_cancel`, `best_bid`, `best_ask` (with `INT64_MAX` normalized to 0 when the ask side is empty, lines 129–134), `mid` (0 if either side is empty), `drain_trades`, `snapshot(depth=10)`, `vwap`, `trade_count`, and — added to close the matcher seam described below — `uncross` (one uniform-price call-auction clearing of the resting book) together with `set_trading_state/trading_state` over the session states of `Auction.tla` (lines 163–189). Fills are captured by a per-symbol `TradeCollector` that buffers every trade on the book and clears on `drain_trades`; Python attributes each fill to agents by `buyer_id/seller_id`. The enums `Side`, `OrderType` (`Limit/Market/IOC/FOK/PostOnly`), `TimeInForce` (`GTC/GTD/DAY`), and `RejectReason`, plus read-only `Trade`, `PriceLevel`, and `MarketDataSnapshot` classes, are exposed alongside.

Python agents drive the verified engine entirely through this surface: each tick they emit orders and cancels, the scheduler submits them via `submit_order/submit_cancel`, and because the engine is synchronous, the scheduler drains the resulting fills inline and routes them back to the agents. The agents never touch C++ state directly; the bridge is the only channel, which is what lets the verified matcher remain the sole adjudicator of trades.

An earlier revision of this harness carried a methodological seam here: the bridge did not expose `uncross`, so the frequent batch (call) auction could not be run through the engine at all and was instead re-implemented in Python (`discover_uncross_price`), replicating `OrderBook::discoverUncrossPrice` by source inspection. That seam is now closed. The bridge exposes `uncross`, `set_trading_state`, and `trading_state`; the batch arm (`BatchAuctionScheduler`) parks the book in `AuctionOpen` — where orders accumulate instead of matching on arrival — collects the window’s orders, and clears through the engine’s own `uncross`, the code path `Auction.tla` proves `SingleClearingPrice` over. Both arms of

Experiment 4 therefore run on the same verified matcher, and the batch arm’s fills are engine fills, drained through the same `TradeCollector` as continuous fills.

A.4 The tie-break divergence

Closing the seam produced a finding in its own right. The retired Python re-implementation had been checked against the C++ source line by line, and an earlier draft of this paper described it as replicating the rule “faithfully” — a claim we noted was source-verified rather than results-verified. When the two implementations could finally be run against the same book, they disagreed. On a book where two prices tie on maximum executable volume (bids 102×50 and 101×30 against asks 99×40 and 100×60 , so 80 units clear at either 100 or 101), the Python rule clears at 100 and the verified engine clears at 101 — a divergence inside the tie-break cascade, invisible to any volume-based check and material to every metric that depends on execution price, rent above all. The divergence is pinned in the test suite (`tests/test_batch_auction.py::test_python_matcher_divergence_pinned`) so it cannot regress silently, and `discover_uncross_price` is retained only as a cross-check. The episode is a compact argument for this paper’s central methodological claim: a re-implementation that is source-verified can still differ from the verified artifact the first time it is results-verified.

The divergence is causally localized to the uncross path, and the experiment contains its own control. A natural objection to the preceding paragraph is that two runs of a stochastic simulation differ for many reasons, so a changed result need not indict the matcher at all. The design answers this objection without requiring us to argue it. Experiment 4’s long grid comprises seven batch arms and one continuous arm, and the continuous arm does not call `uncross` — it matches on arrival, exactly as Experiments 1–3 do. When the grid is re-run with the batch clearing rerouted from the retired Python re-implementation to the engine, the continuous arm returns 578.1 [449, 709], **byte-identical** to its pre-port value, while HFT rent in three batch cells moves across zero in both significance and sign and the maker’s PnL delta changes sign or significance in five (§5.2). Seeds, fundamental paths, agent realizations, run length and every environment parameter are held fixed across the two runs; the single thing that differs is which implementation clears the batch. A control arm that is bit-for-bit unchanged therefore rules out seed variation, drift in the agent code, and environment differences as explanations, and confines the effect to the code path that was substituted. This is a stronger form of evidence than the tie-break counterexample above: the counterexample shows the two rules *can* disagree on a constructed book, whereas the byte-identical control shows the disagreement *did* propagate into published economic quantities and nothing else plausibly did. We draw the general methodological lesson in §5.2, but state the identification here, in the methods, because it is a property of the experimental design rather than an interpretation of the results.

The same episode is a robustness test of an economic assumption, and part of the batch result fails it. The engine clearing at 101 is correct relative to its own specification — maximize executable volume, then minimize imbalance, then minimize distance to a reference price (the last clearing price, or the first order’s price before any trade), then market pressure — not relative to any particular venue, and venues order these steps differently. The tie-break is therefore an assumption about how a batch market prices a volume tie, and the port is a two-rule comparison of it. On the 3,000-tick grid the rule moves nothing: every batch cell keeps the sign and significance of both the maker’s delta and HFT rent. On the 50,000-tick grid it moves most of the batch arm. Under the retired rule the maker still loses at every interval from 5 to 100 ticks, recovering 40–68% of its continuous-clearing loss at intervals 5–25; under the engine’s rule it breaks even or gains, significantly so at intervals 10 (+445) and 25 (+322). What survives both rules is the comparison with continuous clearing: at every interval of five ticks or more, the maker does better batched than continuous. How much better, whether it ends up gaining, and whether batching removes the rent all depend on the tie-break.

Appendix B. Replication

All experiment code is in `bcs_research/`: agents in `agents/`, the latency and batch-auction schedulers in `simulation/`, metrics in `metrics/`, experiment drivers in `experiments/`, analysis scripts in `analysis/`, and the pybind11 bridge in `bindings/engine_bindings.cpp`. The engine’s TLA+ specifications are in `spec/` (notably `MatchingEngine.tla` and `Auction.tla`), with the TLC state-exploration record in `docs/Verification.md`. Results are written as JSON under `bcs_research/results/experiments/`, at $n = 100$ seeds and bootstrap seed 0 unless stated.

Drivers are under `experiments/` and results under `results/experiments/` unless another directory is named.

- **Experiment 1** (§4.3, Table D1): `exp1_hardening.py` → `exp1_hardening.json`; operating point CFG in `exp1_primary.py`.
- **Experiment 2** (§4.3): `exp2_latency_sweep.py` → `exp2_latency_sweep.json`.
- **Experiment 3** (§4.3, Table D2): `exp3_hft_count.py` → `exp3_hft_count.json`. At fixed depth: `main(cfg={"thin_liquidity": False})` → `exp3_hft_count_fixed_depth.json`.
- **Experiment 4** (§4.3): `exp4_batch_auction.py` → `exp4_batch_auction.json`; at 50,000 ticks, `main(cfg={"duration_us": 50_000_000})` → `exp4_batch_long.json`; maker-latency arms, `main(intervals=[1], cfg={"mm_latency_us": 0})`, the same at 1000, and `intervals=[5]` at 7000 → `exp4_batch_mm_lat0_iv1.json`, `exp4_batch_mm_lat1_iv1.json`, `exp4_batch_mm_lat7_iv5.json`.
- **Calibrated re-runs** (§4.4): `run_calibrated.py` → `*_calibrated.json`, `*_calibrated_hftqty50.json`; competing makers, `exp4_competing_makers.py` → `exp4_competing_makers.json`.
- **Incidence surface** (§4.1): `exp6_incidence_surface.py` → `exp6_incidence_surface.json`; leg decomposition, `exp7c_leg_decomposition.py` → `exp7c_leg_decomposition.json`.
- **Maker-count surface** (§4.2): `exp7_3d_surface.py`, run as `main(n_seeds=100, out_name="exp7_3d_surface_n100.json")`.
- **Per-maker decomposition** (§4.2): `exp7b_per_maker_pnl.py` → `exp7b_per_maker_pnl.json`; latency permutations and values, `exp7d_latency_permutation.py` with `perm2`, `perm3` or values → `exp7d_latency_permutation*.json`.
- **Elastic demand** (§4.2): `exp8_elastic_demand.py` → `exp8_elastic_demand.json`.
- **Adverse-sensitivity check** (§5.2): `exp6b_sensitivity_adverse.py` → `exp6b_sensitivity_adverse.json`.
- **Calibration** (§3.6, §3.7, §4.4): `calibration/targets.py` and `calibration/fit.py` → `targets_btcsdt.json`, `fit_btcsdt.json`; `calibration/depth_replenishment.py` → `depth_replenishment.json`; `calibration/trade_book_join.py` → `trade_book_join.json`, all under `results/calibration/`.
- **Hawkes diagnostic** (Appendix C): `analysis/exp5_redesigned.py` → `exp5_redesigned.json`.

Operating point and outputs. Per-seed outputs (`mm_pnl`, `nt_pnl`, `hft_rent` and `_per_sec`, `zero_sum_residual`, `liquidity_gaps`, `kyle_lambda`, `spread_ticks`, `n_trades`, `mm_max_half_spread`, `mm_min_qty`, `hft_snipes`, `hft_fills`) are written to `results/experiments/exp1_primary.json`; the hardened $n = 100$ results reported in §4.3 come from `exp1_hardening.py` and are written to `exp1_hardening.json`. The code names of the parameters this paper describes in words: adverse sensitivity `adverse_sensitivity`, spread decay `spread_decay`, spread cap `spread_cap_mult`, minimum quote size `min_quote_qty`, quote size `quote_qty`, snipe size `hft_qty`, per-tick arrival probability `lambda_per_tick`, HFT latency `hft_latency_us`, the

depth-thinning switch `thin_liquidity`, and the zero-latency-leader arms `zero_fast`.

Metric implementations. Welfare decomposition (`metrics/welfare_decomposition.py`). `decompose_welfare(result, mm_id, nt_ids, hft_ids, mark)` computes per-agent PnL as `cash + inventory * mark / PRICE_PRECISION`, marked at `mark_price` (the last two-sided mid, else the last V). It reports `mm_pnl`, `noise_trader_pnl` (summed over `nt_ids`), `hft_pnl`, `hft_rent` (equal to `hft_pnl` — the rent the arms race extracts), per-agent and per-second variants, and `zero_sum_residual = mm_pnl + nt_pnl + hft_pnl`.

Lagged Kyle’s lambda (`metrics/baseline_metrics.py`, `compute_metrics`). `kyle_lambda` is the OLS slope (`np.polyfit` degree 1) of the *forward-window* mid change `mid[t+k] - mid[t]` (default `k = kyle_horizon_ticks = 5`) regressed on `signed_volume` at `t`, over ticks with `mid > 0`; it requires at least 3 points and `std(x) > 0`, and `kyle_lambda_r2` is the squared correlation. The forward lag, rather than the contemporaneous tick, is deliberate and documented (lines 64–70): a latency-disadvantaged maker quotes stale, so a same-tick regression would score its delayed catch-up requote as *negative* impact; measuring over the next `k` ticks captures the catch-up as the *positive* permanent impact it actually is, removing the stale-quote sign artifact.

Liquidity-gap detector (`metrics/liquidity_gap_detector.py`). As defined in §3.3, `detect_liquidity_gaps` returns the bounded, recovering, endogenous one-sided episodes; `count_liquidity_gaps` returns the count. This is the secondary endogenous-instability metric, matched to the thinning mechanism (gaps, not mid crashes), and each event records `start_t`, `recovery_t`, `duration_us`, `gap_ticks`, `side_missing`, and `v_move`.

Percentile bootstrap CI (`metrics/bootstrap.py`). `bootstrap_ci(values, n_bootstrap=2000, ci=0.95, seed=0)` is a nonparametric percentile bootstrap for the mean: it draws (`n_bootstrap`, `n`) resamples with replacement via `np.random.default_rng(seed)`, takes per-row means, and returns the (seed-independent) point estimate `mean`, the lower/upper bounds at the (`alpha/2`, `1 - alpha/2`) percentiles, `std` (`ddof=0`), and `n`. A single value yields a zero-width CI. `bootstrap_ci_table` builds a `{metric_key: ci}` table from per-seed row dicts. The resampling RNG is explicitly seeded for reproducibility. The `n_bootstrap=2000` resample count and percentile method are defaults of the bootstrap module (source-level), not fields stored in the results JSON, which records only the bounds and `n`. We note that `exp1_primary.py` itself does not call `bootstrap_ci` — it aggregates per-seed rows with `statistics.fmean` — and the bootstrap module is shared infrastructure used from Experiment 2 onward; the `n=100` CIs the paper reports for Experiment 1 come from the hardened run, which does use the bootstrap.

The baseline gap figure. An $n = 20$ run once reported a no-HFT baseline of 3.45 gaps against the $n = 100$ figure of 2.84; raising every experiment to $n = 100$ resolved the discrepancy, confirming it was sample size.

Appendix C. The Hawkes endogeneity diagnostic: a placebo-validated null

Experiment 5 asks whether the liquidity gaps of Experiments 1–3 are preceded by rising order-flow self-excitation — whether they are reflexive cascades rather than direct consequences of adverse-selection widening. We estimate the branching ratio of a self-exciting Hawkes process on simulated order flow (Filimonov and Sornette 2012) in windows labelled pre-gap or normal, on both the trade and quote-update processes, with a placebo arm whose gap labels carry no information. The windowing, labelling, and placebo construction are in the replication materials (Appendix B).

Arm / process	pre mean $n(t)$	normal mean $n(t)$	p (one-sided)	rank-biserial	$n_{\text{pre}} / n_{\text{normal}}$
Treatment / trades	0.0162	0.0167	0.258	0.012	1,416 / 3,614
Treatment / quotes	≈ 0	≈ 0	0.937	-0.029	1,260 / 3,206
Placebo / trades	0.0111	0.0168	0.895	-0.038	373 / 8,614
Placebo / quotes	≈ 0	≈ 0	0.491	0.001	335 / 7,707

The result is a null, and the placebo confirms the pipeline does not manufacture endogeneity where none exists. Pre-gap and normal branching ratios are indistinguishable on trade arrivals ($p = 0.258$, effect size 0.012), and the quote-update process shows essentially zero self-excitation anywhere — mechanically sensible, since the maker’s requotes track an exogenous random walk. At this operating point the gaps are therefore not preceded by rising self-excitation: they arise from the adverse-selection widening mechanism directly. This is a negative result about the simulated mechanism and is stated as such. It is not evidence against reflexivity in real markets, where self-excitation is well documented; this market’s arrival processes are close to Poisson by construction, with branching ≈ 0.017 on trades, some $50\times$ below criticality, leaving little endogeneity to detect. What the experiment contributes is a validated instrument, ready for a real-data leg that awaits a depth-depletion gap definition suited to a book that never empties (§5.3). No figure accompanies this result.

Appendix D. Supporting tables

Table D1. Experiment 1 welfare and microstructure estimates at the operating point ($n = 100$ seeds; point estimate with 95% bootstrap CI).

Quantity	Treatment	No-HFT baseline	Treatment – baseline
HFT rent	34.70 [26.25, 44.80]	—	—
MM PnL	72.78 [71.05, 74.38]	101.27 [93.28, 110.72]	-28.50 [-38.31, -20.35]
NT PnL	-107.48 [-116.84, -99.26]	-101.27 [-110.72, -93.28]	-6.21 [-7.87, -4.61]
Zero-sum residual	2.78e-12 [-3.34e-12, 9.11e-12]	-5.39e-12 [-2.16e-11, 1.05e-11]	—
Liquidity gaps	9.70 [9.15, 10.27]	2.84 [2.47, 3.26]	6.86 [6.20, 7.52]
Lagged Kyle’s λ	1.372 [1.361, 1.383]	0.200 [0.184, 0.216]	—

Table D2. Experiment 3: per-HFT and total rent, fill rate and liquidity gaps by HFT count k at the operating point ($n = 100$ seeds; point estimate with 95% bootstrap CI).

k	Per-HFT rent	Total rent	Fill rate	Liquidity gaps
0	—	—	—	2.84 [2.47, 3.26]
1	34.82 [26.40, 44.84]	34.82 [26.40, 44.84]	0.2886 [0.2863, 0.2909]	5.27 [4.91, 5.65]
2	17.46 [13.28, 22.43]	34.93 [26.55, 44.87]	0.1448 [0.1437, 0.1460]	7.79 [7.31, 8.29]

k	Per-HFT rent	Total rent	Fill rate	Liquidity gaps
3	11.57 [8.75, 14.93]	34.70 [26.25, 44.80]	0.0964 [0.0957, 0.0971]	9.70 [9.15, 10.27]
5	6.77 [5.07, 8.81]	33.83 [25.36, 44.03]	0.0578 [0.0574, 0.0583]	14.92 [14.27, 15.61]
8	4.31 [3.23, 5.60]	34.50 [25.84, 44.76]	0.0362 [0.0359, 0.0365]	21.09 [20.20, 21.99]
13	2.65 [1.97, 3.45]	34.43 [25.63, 44.87]	0.0223 [0.0222, 0.0225]	32.79 [31.73, 33.91]
21	1.60 [1.21, 2.07]	33.70 [25.39, 43.45]	0.0139 [0.0137, 0.0140]	47.47 [46.22, 48.85]

References

- Aquilina, M., Budish, E., & O'Neill, P. (2022). Quantifying the High-Frequency Trading “Arms Race”. *Quarterly Journal of Economics*, 137(1), 493–564.
- Baldauf, M., & Mollner, J. (2020). High-Frequency Trading and Market Performance. *Journal of Finance*, 75(3), 1495–1526.
- Biais, B., Foucault, T., & Moinas, S. (2015). Equilibrium Fast Trading. *Journal of Financial Economics*, 116(2), 292–313.
- Brogaard, J., Hendershott, T., & Riordan, R. (2014). High-Frequency Trading and Price Discovery. *Review of Financial Studies*, 27(8), 2267–2306.
- Budish, E., Cramton, P., & Shim, J. (2015). The High-Frequency Trading Arms Race: Frequent Batch Auctions as a Market Design Response. *Quarterly Journal of Economics*, 130(4), 1547–1621.
- Carlin, B. I., Lobo, M. S., & Viswanathan, S. (2007). Episodic Liquidity Crises: Cooperative and Predatory Trading. *Journal of Finance*, 62(5), 2235–2274.
- Copeland, T. E., & Galai, D. (1983). Information Effects on the Bid-Ask Spread. *Journal of Finance*, 38(5), 1457–1469.
- Du, S., & Zhu, H. (2017). What Is the Optimal Trading Frequency in Financial Markets? *Review of Economic Studies*, 84(4), 1606–1651.
- Farmer, J. D., & Foley, D. (2009). The Economy Needs Agent-Based Modelling. *Nature*, 460(7256), 685–686.
- Filimonov, V., & Sornette, D. (2012). Quantifying Reflexivity in Financial Markets: Toward a Prediction of Flash Crashes. *Physical Review E*, 85(5), 056108.
- Foucault, T., Kozhan, R., & Tham, W. W. (2017). Toxic Arbitrage. *Review of Financial Studies*, 30(4), 1053–1094.
- Glosten, L. R., & Milgrom, P. R. (1985). Bid, Ask and Transaction Prices in a Specialist Market with Heterogeneously Informed Traders. *Journal of Financial Economics*, 14(1), 71–100.
- Hoffmann, P. (2014). A Dynamic Limit Order Market with Fast and Slow Traders. *Journal of Financial Economics*, 113(1), 156–169.
- Kirilenko, A., Kyle, A. S., Samadi, M., & Tuzun, T. (2017). The Flash Crash: High-Frequency Trading in an Electronic Market. *Journal of Finance*, 72(3), 967–998.
- Kyle, A. S. (1985). Continuous Auctions and Insider Trading. *Econometrica*, 53(6), 1315–1335.

- LeBaron, B. (2006). Agent-Based Computational Finance. In *Handbook of Computational Economics*, Vol. 2 (pp. 1187–1233).
- Menkveld, A. J. (2013). High Frequency Trading and the New Market Makers. *Journal of Financial Markets*, 16(4), 712–740.
- Menkveld, A. J. (2016). The Economics of High-Frequency Trading: Taking Stock. *Annual Review of Financial Economics*, 8, 1–24.
- Menkveld, A. J., & Zoican, M. A. (2017). Need for Speed? Exchange Latency and Liquidity. *Review of Financial Studies*, 30(4), 1188–1228.
- Shkilko, A., & Sokolov, K. (2020). Every Cloud Has a Silver Lining: Fast Trading, Microwave Connectivity, and Trading Costs. *Journal of Finance*, 75(6), 2899–2927.
- Wah, E., & Wellman, M. P. (2013). Latency Arbitrage, Market Fragmentation, and Efficiency: A Tactical Market Model. In *Proceedings of the 14th ACM Conference on Electronic Commerce (EC)*, 855–872.