

High-Performance Order Matching Engine — Project Overview

21,600+ lines of C++20 | 30 headers | 8 source files | 30 test files | 28 test executables

A production-grade, institutional-architecture matching engine built for sub-100ns processing latency and horizontal scalability.

1. Executive Summary


This is a C++20 low-latency order matching engine with institutional-grade infrastructure modeled after CME/NYSE exchange architecture. It implements $O(1)$ price-level lookup via `FlatPriceMap`, lock-free MPSC queues, thread-per-symbol horizontal scaling, CRC-32 journaling with deterministic replay, FIX 4.4 session management, TLA+-verified concurrency, primary-backup high availability with epoch-based leader fencing, and a complete operational stack including config management, webhook alerting, Prometheus metrics, and Docker deployment.

Codebase Statistics

Metric	Count
Total C++ LOC	21,630+
Header files (<code>include/</code>)	30
Source files (<code>src/</code>)	8
Test files (<code>tests/</code>)	30
Test executables	28
Individual test cases	220+
TLA+ specifications	4 (750k+ states verified)
Documentation files	6

2. Core Engine Architecture

System Topology



ObjectPool
Journal

Intrusive linked lists
CRC-32 WAL

Order Types Supported

- **Standard:** Limit, Market
- **Time-In-Force:** IOC (Immediate-or-Cancel), FOK (Fill-or-Kill), DAY, GTD (Good-Til-Date)
- **Conditional:** Stop, StopLimit, TrailingStop
- **Institutional:** Pegged, Iceberg (hidden quantity), Hidden, PostOnly
- **Match Algorithms:** Price-Time FIFO and Pro-Rata allocation

Core Data Structures

Component	Purpose	Complexity
FlatPriceMap	Price-level lookup by tick index	O(1) insert/lookup
FlatHashMap	Robin-Hood open-addressing hash map	O(1) amortized
IntrusiveList	Doubly-linked order queue per price level	O(1) insert/remove
ObjectPool	Pre-allocated slab allocator for Order nodes	O(1) alloc/dealloc
RingBuffer	Cache-line aligned lock-free ring buffer	O(1) push/pop
MpscQueue	Multi-producer, single-consumer lock-free queue	O(1) push/pop

Data Flow

1. **Ingress:** Order arrives via TCP Gateway (binary framing or FIX 4.4)
2. **Rate Limiting:** Token-bucket per-participant throttling at ingress
3. **Queue Routing:** `hash(symbolId) % numThreads` → lock-free `MpscQueue`
4. **Validation:** Circuit Breaker check, OTR enforcement, risk limit verification
5. **Matching:** Price-time priority against resting book (`FlatPriceMap`)
6. **Execution:** Trade records generated, callbacks invoked
7. **Persistence:** CRC-32 journal entry written (WAL)
8. **Analytics:** VWAP, TWAP, and OTR stats updated
9. **Market Data:** L2 updates published via shared memory (`shm_open`)

3. Concurrency & Networking

Thread-Per-Symbol Partitioning

N worker threads with independent lock-free `MpscQueue` ring buffers. Orders are deterministically routed by `hash(symbolId) % numThreads` , eliminating cross-thread contention on the hot path.

TCP Gateway (Binary Protocol)

- Non-blocking I/O via `epoll` (Linux) / `kqueue` (macOS)
- `GatewayProtocol.h` — 16-byte fixed header with V1/V2 payload evolution
- Forward and backward compatibility via version negotiation
- `GatewayResponse` ack frames with numeric reject reasons

FIX 4.4 Gateway

- Full FIX 4.4 session management (`FixSession.h`)
- `FixFramer` for message reassembly across TCP fragments
- `FIXParser` — zero-copy tag-value parser

- `FixTcpGateway` — production-grade FIX acceptor
- `TransactTime` enforcement and `OrdRejReason` numeric mapping

Shared Memory IPC

- `MarketDataPublisher` using POSIX `shm_open` with versioned `ShmHeader`
- Magic/version/entrySize fields for prefix-compatible rolling upgrades
- Low-overhead L2 market data distribution to co-located consumers

4. Regulatory & Risk Management

Feature	Implementation
Self-Match Prevention (SMP)	Cancel-Taker strategy — prevents wash trading
Circuit Breakers	Configurable % price band halts (volatility protection)
OTR Monitoring	Real-time Order-to-Trade ratio tracking per participant
Kill Switch	Instant cancellation of all orders per participant across all symbols
Pre-Trade Risk Limits	Max order size, notional value, and position limits per participant
Rate Limiting	Token-bucket throttling at ingress (configurable rate + burst)
Queue Backpressure	Rejects orders when queue exceeds configurable threshold (default 80%)

5. Reliability & High Availability

CRC-32 Journaling

- Write-ahead log with atomic CRC-32 integrity on every entry
- Crash recovery: replay stops at first corrupted/truncated record
- Atomic checkpoint: snapshots active book state, rewrites journal
- Deterministic replay via virtual clock (`setExpiryClock(ClockFn)`)

Primary-Backup Replication

Modeled after CME/NYSE exchange HA architecture:

- `ReplicationCoordinator` — TCP log-shipping from primary to backup
- `HeartbeatMonitor` — configurable failure detection timeout
- `LeaderLease` — epoch-based fencing prevents split-brain
- `JournalFollower` — single-host warm standby with `promote()` API
- Automatic backup promotion on primary failure detection

Deterministic Sequencing

- Monotonic `sequenceNumber` on all `Trade` , `OrderUpdate` , and `MarketDataUpdate` events
- Gap detection for downstream consumers
- GTD virtual clock seam for cross-day replay — DAY/GTD expirations are journaled and replayed identically

6. Observability & Operations

Structured Logging

- Pluggable `StructuredSink` API
- `NullSink` — zero-cost in production (compiled out)
- `JsonStderrSink` — development diagnostics

- `CapturingSink` — test assertions on log output

Prometheus Metrics

- `MetricsRegistry` with atomic Counters, Gauges, Histograms
- `/prometheus` text-exposition endpoint for standard Prometheus scraping
- Tracks: order flow, rejects, queue depth, latency histograms, journal health, gateway sessions

Webhook Alerting

- `AlertDispatcher` — background thread delivery
- Targets: Slack, PagerDuty, or generic HTTP webhooks
- Configurable severity filtering (Info, Warning, Critical)

Configuration Management

- `Config` key-value loader with env var override (`OB_` prefix)
- Type-safe getters (`getInt` , `getString` , `getBool`)
- Hot-reload support for runtime-adjustable parameters
- Example config: `config/engine.conf.example`

Admin HTTP Server (Port 8080)

Endpoint	Purpose
GET <code>/health</code>	K8s liveness probe
GET <code>/metrics</code>	Internal counters (JSON)
GET <code>/prometheus</code>	Prometheus text exposition
GET <code>/book?symbolId=0</code>	L2 order book snapshot
GET <code>/otr?participantId=1</code>	Order-to-trade ratio

Docker Deployment

- Multi-stage `Dockerfile` for minimal production image
- `docker-compose.yml` for primary-backup topology
- Health checks and journal volume persistence

7. Performance Benchmarks

Measured locally using hardware timing on Apple Silicon M-series. All numbers are from `./bin/ManualBenchmark` and should be treated as machine-specific results, not a portable latency guarantee.

Note: These are **per-order processing latencies** — how long each operation takes once it is dequeued and executing on the matching thread. They do **not** include queuing delay.

Lean Mode (`OB_LEAN_MODE`)

Disables risk checks, OTR tracking, and event notifications for raw matching speed.

Operation	Avg	P50	P99	Throughput
Add	39.8 ns	42 ns	42 ns	25.2M ops/s
Match	70.1 ns	83 ns	125 ns	14.3M ops/s
Cancel	32.4 ns	0 ns	208 ns	30.8M ops/s

Full Mode (Enterprise Default)

Includes SMP, circuit breakers, OTR, risk limits, and full event routing.

Operation	Avg	P50	P99	Throughput
Add	520.3 ns	541 ns	1167 ns	1.9M ops/s
Match	109.2 ns	84 ns	250 ns	9.1M ops/s
Cancel	1092.3 ns	875 ns	6208 ns	0.9M ops/s

Multi-Threaded Stress Test (4 threads)

Scenario	Orders	Trades	Result
Concurrent Producers	20,000	~6,300	✓
Mixed Workload	20,000	~5,800	✓
Kill Switch Under Load	11,000	—	✓
Multi-Symbol (4 symbols)	12,000	~5,900	✓
High-Throughput Burst	50,000	~19,500	✓ 5M ops/s

End-to-End Latency (ingress → completion, including queue delay)

Scenario	Threads	P50	P99	P99.9	Throughput
Single order (drain between submits)	2	167 ns	292 ns	3.58 μs	3.30M accepted ops/s
10K burst (no rate limit)	2	208 ns	30.21 μs	32.13 μs	7.78M accepted ops/s
Rate-limited (5K/s, 10 producers)	2	3.20 μs	57.86 μs	109.06 μs	0.04M accepted ops/s
Rate-limited (10K/s, 20 producers)	4	5.82 μs	52.74 μs	74.24 μs	0.17M accepted ops/s

Tail latency bounding: Per-participant token-bucket rate limiting caps aggregate inbound rate. Queue-depth backpressure (configurable threshold, default 80%) rejects orders before the queue saturates. Together: `max_queue_depth × avg_processing_time` gives a hard ceiling on queuing delay.

8. Verification & Testing

Test Suite — 28 Executables, 220+ Test Cases

Category	Tests	Description
Functional	ManualTest, UnitTests (GTest)	All order types, matching, cancellation, edge cases
Stress	StressTest	5 multi-threaded scenarios, 4 concurrent producers
Property	PropertyTest, JournalReplayPropertyTest, SnapshotConsistencyTest	Randomized invariant checks, replay determinism
Chaos	FaultInjectorTest, PoolExhaustionTest, JournalChaosTest, JournalCorruptionTest, GatewayChaosTest, QueueChaosTest,	10+ fault-injection points, short-writes, bit-flips,

	CheckpointChaosTest, CombinedChaosTest	EAGAIN
Integration	GatewayIntegrationTest, FixTcpGatewayTest, AdminServerEndpointsTest	TCP round-trip, FIX session, HTTP endpoints
Protocol	GatewayProtocolTest, MarketDataSchemaTest, MdFeedSchemaCompatTest, RejectReasonContractTest, SubmitResultTest	Binary wire format evolution, ShmHeader compat
Recovery	JournalCrashTest, JournalFollowerTest, GTDReplayTest	Crash recovery, log-ship convergence, virtual-clock replay
HA & Ops	ReplicationProtocolTest, ConfigTest, AlertDispatcherTest	Primary-backup replication, config loader, webhook alerts
Benchmark	BenchmarkRegression (GTest)	P99 latency regression gates

Formal Verification (TLA+)

Specification	States Explored	Invariants
MpscQueue.tla	~250K	Lock-free ring buffer linearizability
EngineConsumer.tla	~200K	Worker loop shutdown safety
Snapshot.tla	~150K	Read/write mutex prevents torn snapshots
SnapshotLocked.tla	~150K	Locked variant correctness
Total	~750K	All invariants hold

Chaos Engineering & Fault Injection

`FaultInjector` singleton with 10+ injection points:

- Journal: short-writes, bit-flips, fsync failures
- Memory: pool exhaustion, allocation failures
- Network: gateway fragmentation, EAGAIN injection
- Queue: spurious empty returns, backpressure simulation
- **Zero-cost in production** (`OB_ENABLE_FAULT_INJECTION` compiled out)

Coverage-Guided Fuzzing

- `libFuzzer` harness for protocol parsing and order flow
- `fuzz/` directory with dedicated build configuration

CI/CD Pipeline

GitHub Actions workflow runs on every push/PR:

- **macOS Release:** Full test suite
- **Ubuntu Debug + ASan/UBSan:** Memory and undefined behavior sanitizers
- **Ubuntu ThreadSanitizer:** Data race detection
- E2E benchmark smoke test with artifact upload

9. Capacity Planning

Memory Sizing

Scenario	Live Orders	Symbols	Memory
Small (dev/test)	10K	4	~256 MB
Medium (prop desk)	100K	50	~512 MB
Large (exchange)	1M	500	~2 GB
Ultra (HFT venue)	5M	2000	~8 GB

CPU Scaling

Cores	Workers	Throughput (lean)	Throughput (full)
4	1	25M ops/s	1.9M ops/s
8	4	80M ops/s	7M ops/s
16	8	150M ops/s	14M ops/s
32	16	280M ops/s	25M ops/s

Disk & Network

- **Journal:** ~80-100 bytes/entry. 1M orders/day ≈ 100MB/day
- **Recommended:** NVMe SSD with ≥100K IOPS (production)
- **Network:** 10 GbE for production; 25 GbE + kernel bypass for HFT

10. Deployment

Build & Run

```
# Build with all enterprise features
./build.sh

# Build for sub-100ns lean mode
./build.sh --lean

# Run the engine
./bin/OrderEngine --threads 4 --port 8080 --symbols 4

# Run full test suite
ctest --test-dir build --output-on-failure -L project
```

Docker

```
docker-compose up -d engine-primary
curl -sf http://localhost:8080/health
```

Configuration Precedence

1. Environment variable `OB_<KEY>` (highest priority)
2. Config file value (`config/engine.conf.example`)

11. File Structure

include/	30 header files – all core logic is header-only
├─ OrderBook.h	Core matching engine (373 lines)
├─ FlatPriceMap.h	O(1) price-level lookup (401 lines)
├─ FlatHashMap.h	Robin-Hood hash map (285 lines)
├─ MatchingEngine.h	Thread-per-symbol router (11K bytes)
├─ MpscQueue.h	Lock-free multi-producer queue
├─ Journal.h	CRC-32 WAL + crash recovery
├─ TcpGateway.h	Binary protocol gateway
├─ FIXParser.h	FIX 4.4 tag-value parser
├─ FixSession.h	FIX session state machine
├─ FixTcpGateway.h	FIX TCP acceptor
├─ ReplicationProtocol.h	Primary-backup HA (20K bytes)
├─ Config.h	Key-value config with env override
├─ AlertDispatcher.h	Webhook alerting
├─ Metrics.h	Prometheus metrics registry
├─ StructuredLog.h	Pluggable structured logging
├─ MarketDataPublisher.h	Shared memory market data
├─ GatewayProtocol.h	Binary wire protocol versioning
└─ ...	Types, Utils, LatencyTracker, etc.
src/	8 source files – thin compilation units
tests/	30 test files – 28 executables
spec/	TLA+ formal specifications (4 specs)
config/	Example configuration files
docs/	Runbook, CapacityPlanning, ProductionReadiness
fuzz/	Coverage-guided fuzzing harness
scripts/	Benchmark and build automation

12. Documentation

Document	Purpose
README.md	Feature overview and getting started
Architecture.md	Technical architecture deep-dive
PerformanceWhitepaper.md	Benchmark methodology and analysis
docs/Runbook.md	Operational procedures and incident response
docs/CapacityPlanning.md	Memory, CPU, disk, and network sizing
docs/ProductionReadiness.md	Remaining production checklist
config/engine.conf.example	All configuration keys with defaults
PRODUCTION_ROADMAP.md	Development roadmap and status

13. Remaining Work

The system is architecturally complete. The only remaining gaps require specialized hardware not available in a standard development environment:

Item	Status	Blocker
io_uring zero-copy data path	Seams exist (FixSession/FixFramer)	Linux + io_uring kernel
DPDK kernel bypass	Architecture ready	Linux + supported NIC
Solarflare/Onload	Architecture ready	Solarflare hardware
Wire-to-wire latency measurement	E2E bench exists	Multi-host test rig

Developed for professional quantitative trading systems. C++20 · 21,600+ LOC · 28 test executables · 750K+ TLA+ states verified